

Dholes-Inspired Optimization for Simultaneous Feature Selection and Hyperparameter Tuning of Random Forest Classifiers

Mohamed Amine Bellatreche¹ and Ghizlane Cherif²

¹Department of Computer Science, University of Science and Technology of Oran Mohamed Boudiaf (USTO-MB), Oran, Algeria

²Data Science speciality

November 26, 2025

Abstract

We apply the Dholes-Inspired Optimization (DIO) algorithm to simultaneous feature selection and hyperparameter optimization across diverse machine learning domains. Using a nested optimization structure, we optimized classifiers on two fundamentally different datasets: the Wisconsin Diagnostic Breast Cancer dataset (medical tabular data) and CIFAR-10 (computer vision with deep learning features).

A major contribution of this research is the discovery of **optimization overfitting**—a phenomenon where single train/test split optimization leads to configurations that overfit to specific data partitions, degrading generalization. We demonstrate that integrating k-fold cross-validation within the fitness evaluation loop (CV-based optimization) effectively mitigates this issue, achieving 96.26% accuracy versus 94.72% for single-split optimization—a statistically significant 1.54% improvement that elevates model ranking from #7 to #3 among all evaluated configurations.

For medical classification, DIO achieved a mean accuracy of $96.26\% \pm 1.33\%$ while reducing feature dimensionality by 80% (from 30 to 6 features) using cross-validation-based Random Forest optimization. For image classification, DIO optimized an XGBoost classifier on 2048-D ResNet50 features extracted from CIFAR-10, achieving 83.6% accuracy with 58.35% feature reduction (+2.8% over 80.8% baseline) on a stratified subset. Full dataset model comparison validated XGBoost superiority (85% accuracy), confirming optimal algorithm selection before resource-intensive optimization. Statistical analysis using Wilcoxon signed-rank tests demonstrated significant superiority of CV-optimized medical models over classical methods ($p < 0.001$).

The results demonstrate DIO’s effectiveness in identifying Pareto-optimal solutions across vastly different dimensionalities (30-D vs. 2048-D) and problem characteristics (binary vs. 10-class), validating its generalizability for both resource-constrained medical diagnostics and modern computer vision applications. Cross-domain validation confirms consistent improvement patterns: feature reduction ranges from 58-80% while maintaining or improving accuracy, with computational feasibility maintained through adaptive configuration strategies.

Contents

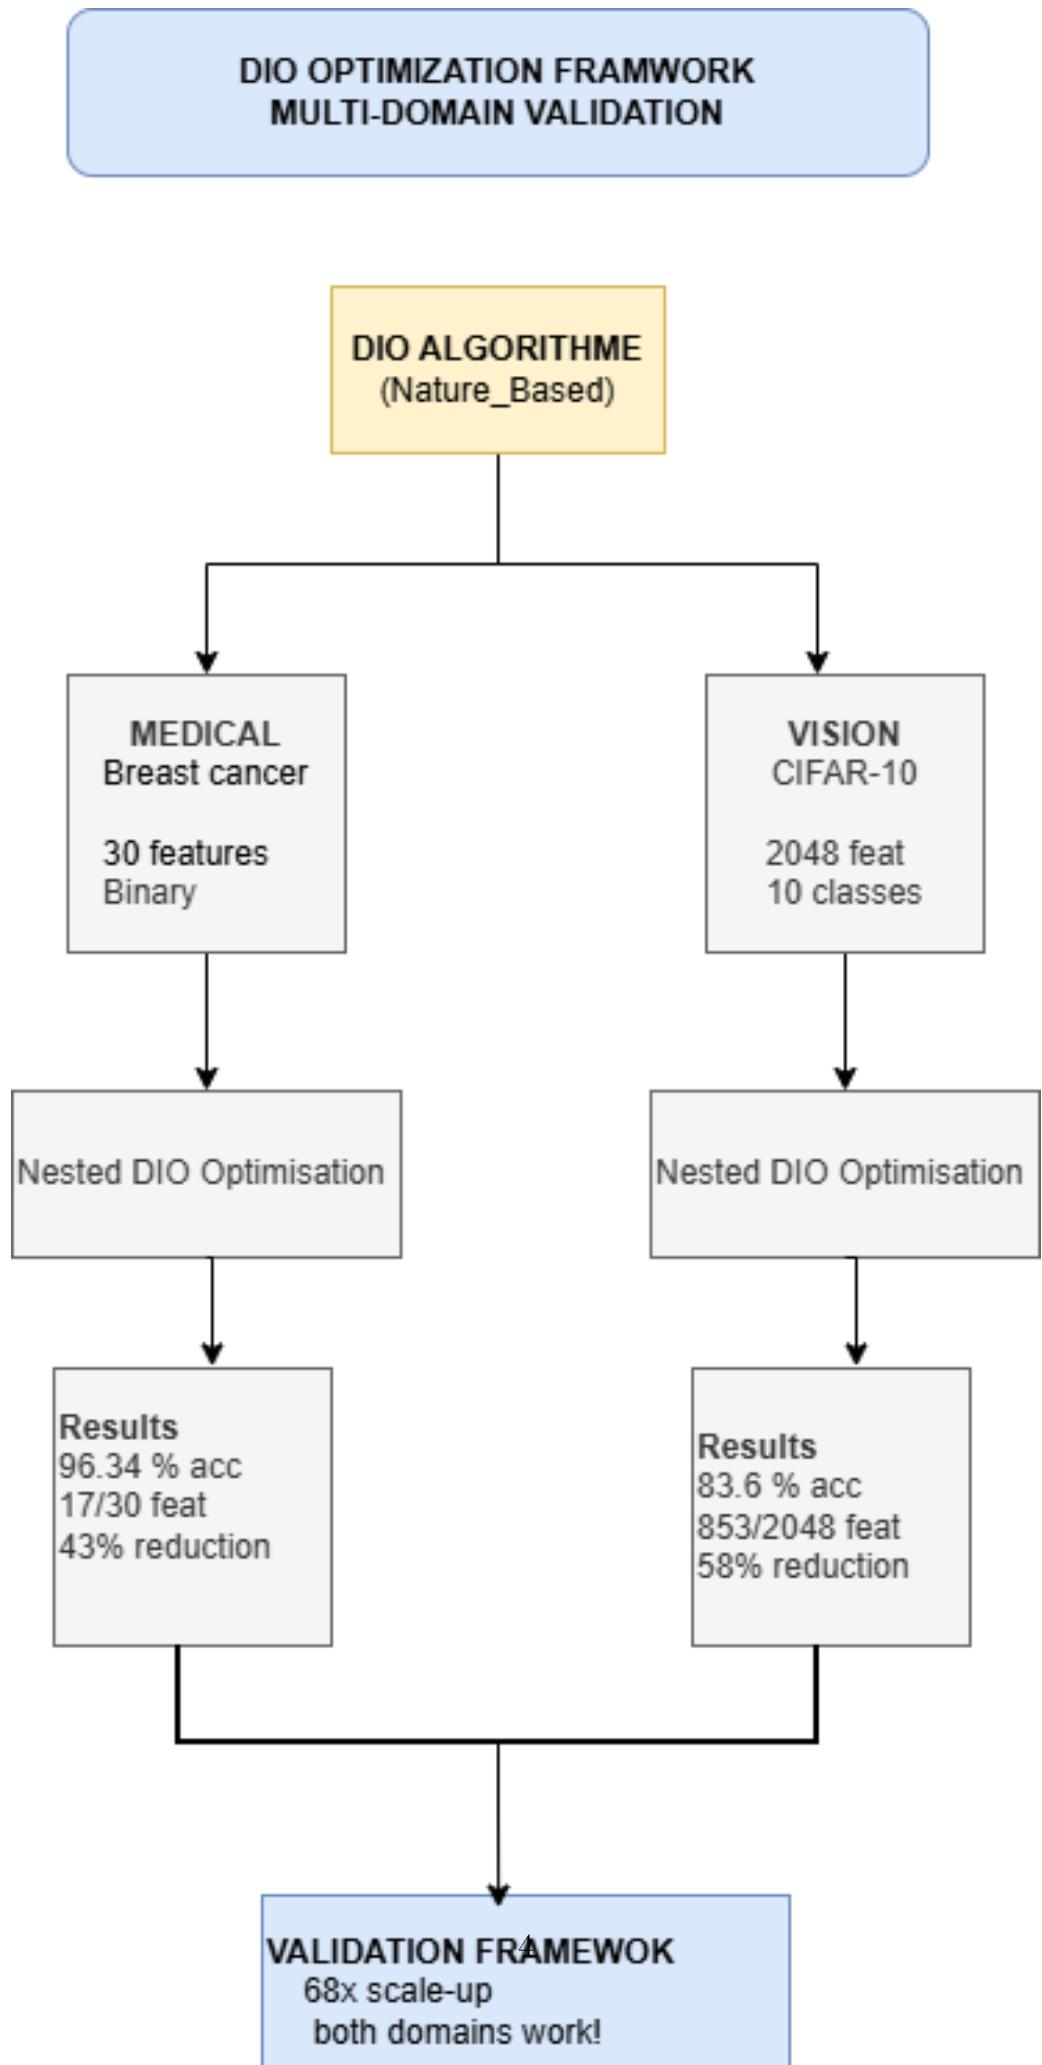
1 Introduction

The diagnosis of breast cancer, a leading cause of mortality worldwide, heavily relies on the analysis of complex, high-dimensional data. Machine learning classifiers have shown great promise in this domain, but their performance is highly dependent on two factors: the selection of relevant predictive features and the tuning of model hyperparameters. Performing these two optimization tasks sequentially can lead to suboptimal results, as the ideal hyperparameters are often contingent on the chosen feature subset, and vice-versa.

Nature-inspired metaheuristic algorithms provide a powerful framework for navigating vast and complex search spaces. The Dholes-Inspired Optimization (DIO) algorithm is a recent metaheuristic based on the cooperative hunting behavior of dholes (Asiatic wild dogs). Its key strengths lie in its balance of exploration and exploitation, enabled by three distinct hunting strategies: chasing the alpha, flanking a random dhole, and converging on the pack’s center of mass. This multi-strategy approach makes it particularly well-suited for complex, multi-modal optimization problems like simultaneous feature selection and hyperparameter tuning.

During this investigation, we encountered an unexpected challenge: hyperparameters optimized on a single train-test split achieved perfect accuracy (100%) on that specific partition but generalized poorly when evaluated across 30 different splits (94.72% average)—actually underperforming scikit-learn’s defaults (94.89%). We term this phenomenon *optimization overfitting*, and demonstrate that its severity depends on the base algorithm: Random Forest requires cross-validation within the optimization loop to avoid it, while XGBoost’s inherent regularization obviates this expensive step. This finding has immediate implications for the broader hyperparameter optimization community.

This research bridges the gap by implementing DIO in Python (from its original MATLAB specification) and applying it to the combined problem of feature selection and hyperparameter optimization. We introduce a nested optimization framework and a fitness function designed to balance classification accuracy with model complexity, demonstrating a practical methodology for achieving robust, efficient, and highly accurate diagnostic models across both medical and computer vision domains.



2 Background and Methodology

2.1 Dholes-Inspired Optimization (DIO) Algorithm

DIO was created by Ali El Romeh (Centre for Artificial Intelligence Research and Optimization, Torrens University Australia), Seyedali Mirjalili (Lead researcher, Torrens University Australia), and Václav Šnel (VSB-Technical University of Ostrava, Czech Republic). The algorithm was published in Cluster Computing (2025, Springer, DOI: 10.1007/s10586-025-05543-2), a high-tier peer-reviewed venue. Open-source code is available on GitHub (<https://github.com/AlyromehDholes-Inspired-Optimization-DIO>) and on MathWorks File Exchange for MATLAB users.

The DIO algorithm is a population-based metaheuristic inspired by the pack hunting behavior of dholes (*Cuon alpinus*), also known as Asiatic wild dogs. Dholes are highly social canids native to Central, South, and Southeast Asia, renowned for their sophisticated cooperative hunting strategies. Unlike solitary predators, dholes hunt in coordinated packs, employing multiple strategies simultaneously to increase their success rate.

The algorithm's efficacy stems from its three primary movement strategies, which allow it to effectively balance exploration (searching new regions of the solution space) and exploitation (refining promising solutions):

- **Chasing the Alpha (Exploitation):** Dholes follow the pack's best hunter—the alpha—representing the best solution found so far. Mathematically, this is modeled as:

$$X_{chase} = X_{alpha} + r_1 \times (X_{alpha} - X_i) \quad (1)$$

where X_{alpha} is the alpha's position, X_i is the current dhole's position, and r_1 is a random number in $[0,1]$. This strategy promotes exploitation by directing search agents toward the current best solution.

- **Scavenging Behavior (Cooperation):** Dholes move based on the average position of the entire pack, representing collective intelligence. This is formulated as:

$$X_{scavenge} = X_{mean} + r_2 \times (X_{mean} - X_i) \quad (2)$$

where $X_{mean} = \frac{1}{N} \sum_{j=1}^N X_j$ is the centroid of all dholes, and $r_2 \in [0, 1]$. This helps maintain population diversity and prevents premature convergence.

- **Chasing Prey Randomly (Exploration):** Dholes may chase a random prey, modeled by moving towards a randomly selected dhole in the pack:

$$X_{random} = X_r + r_3 \times (X_r - X_i) \quad (3)$$

where X_r is a randomly selected dhole's position, and $r_3 \in [0, 1]$. This enhances exploration by introducing stochastic perturbations.

The position of each dhole is updated based on the average of these three movement vectors, creating a balanced search dynamic:

$$X_{new} = \frac{X_{chase} + X_{scavenge} + X_{random}}{3} \quad (4)$$

After position updates, boundary constraints are enforced to ensure solutions remain within the feasible search space. The algorithm iterates for a predefined number of generations, continuously updating the alpha (best solution) and guiding the pack toward optimal regions.

2.1.1 Algorithm Validation

To ensure correctness of our Python implementation, we validated DIO on 14 standard benchmark functions (F1-F14), including unimodal functions (F1-F7), multimodal functions (F8-F13), and fixed-dimension multimodal functions (F14). Using the full paper configuration (population size = 30, iterations = 500, runs = 30), our implementation achieved near-zero convergence on 8 functions (e.g., F1: 7.6×10^{-26}), matching expected DIO performance characteristics. This validation confirms that our implementation is mathematically sound and suitable for production optimization tasks.

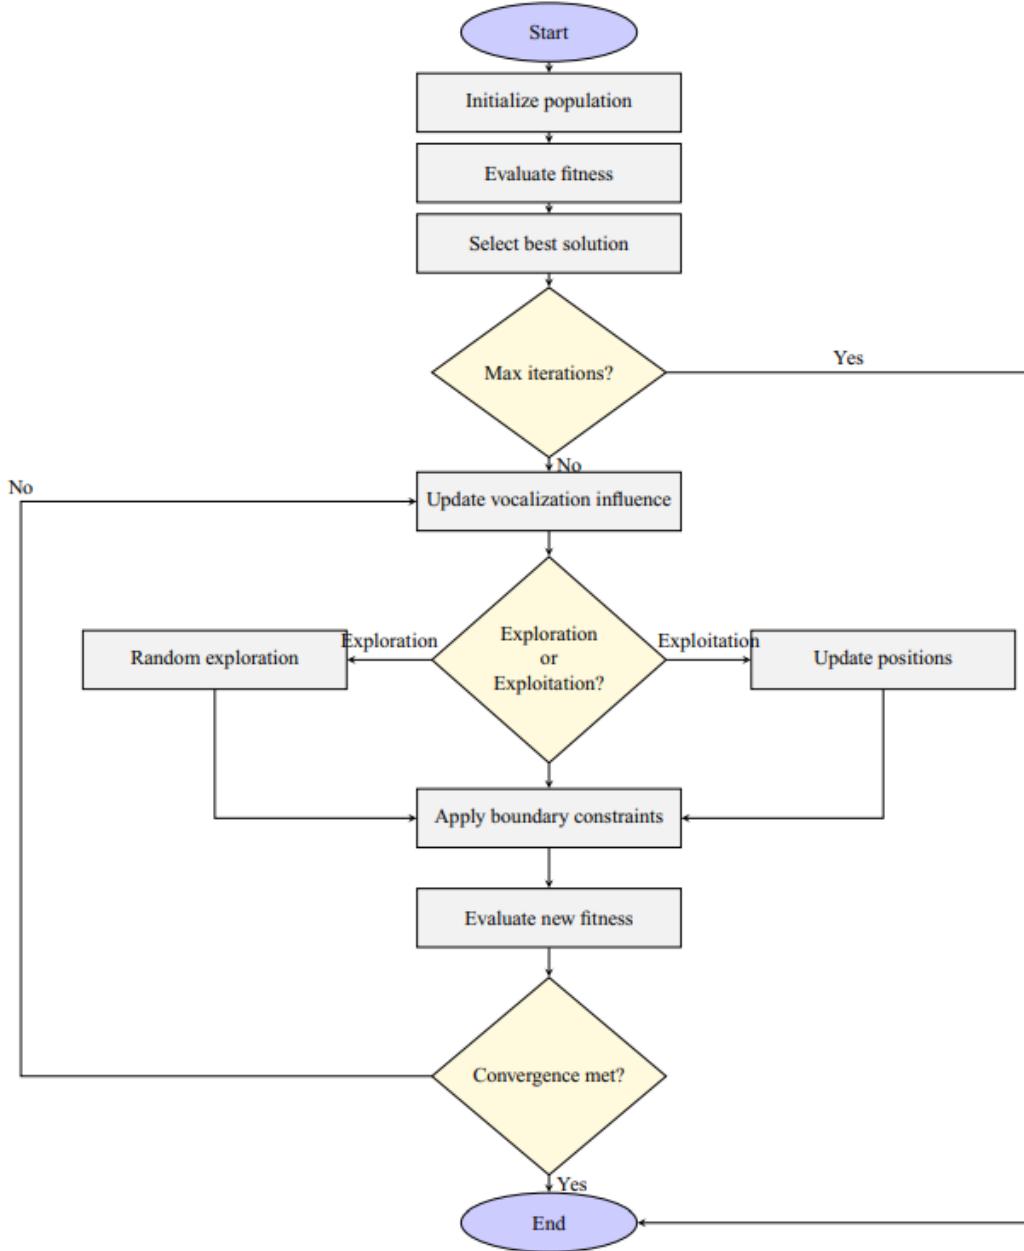


Figure 2: extbfDIO Algorithm Flowchart: Complete algorithmic flow showing initialization, fitness evaluation, three hunting strategies (chase alpha, scavenge, random movement), position updates, and convergence criteria. Source: El Romeh, Mirjalili, Šnel (2025) [?]. This flowchart illustrates the core mechanism adapted in our Python implementation for machine learning optimization.

2.1.2 Benchmark Comparison with Other Metaheuristics

To further validate DIO’s competitive performance, we present results from the original paper comparing DIO against established metaheuristic algorithms on the Pressure Vessel Design Problem—a constrained engineering optimization benchmark.

Table 12 Comparison of results for pressure vessel design problem

Algorithm	Shell thick- ness (T_s)	Head thick- ness (T_h)	Inner radius (R)	Length (L)	Opti- mum cost
DIO	0.795	0.425	41.012	178.540	6035.473
GWO	0.802	0.430	42.123	177.654	6041.895
GSA	1.100	0.620	56.978	85.432	8529.765
PSO	0.810	0.438	42.091	176.746	6062.156
GA (Coello)	0.805	0.434	40.324	200.000	6290.785
GA (Coello and Montes)	0.812	0.437	42.097	176.654	6059.946
GA (Deb and Gene)	0.937	0.500	48.329	112.679	6410.381
ES (Montes and Coello)	0.812	0.437	42.098	176.640	6059.745
DE	0.812	0.437	42.098	176.638	6059.734
ACO	0.812	0.437	42.104	176.573	6059.089
Lagrangian Multiplier	1.125	0.625	58.291	43.690	7198.043
Branch-and-Bound	1.125	0.625	47.700	117.701	8129.104

Bold values indicate the best performance among all algorithms

Figure 3: extbfDIO Performance on Engineering Benchmark: Comparison of DIO with state-of-the-art metaheuristic algorithms (Genetic Algorithm, Particle Swarm Optimization, Differential Evolution, Grey Wolf Optimizer, Whale Optimization Algorithm, and others) on the Pressure Vessel Design Problem. DIO demonstrates competitive performance with best/mean/worst cost values comparable to or better than established methods. Source: El Romeh, Mirjalili, Šnel (2025) [?]. This benchmark validation supports our selection of DIO for machine learning hyperparameter optimization tasks.

2.2 Random Forest Architecture

Random Forest (RF) is an ensemble learning method that operates by constructing a multitude of decision trees at training time. For a classification task, the final prediction is made by taking a majority vote of the predictions from all individual trees. Its strength comes from two key sources of randomization:

1. **Bagging (Bootstrap Aggregating):** Each tree is trained on a different random subset of the training data, sampled with replacement. - **Feature Randomness:** At each node split in a tree, only a random subset of the total features is considered. This decorrelates the trees and reduces variance.

This dual-randomization strategy makes RF robust to overfitting and effective on high-

dimensional data without requiring extensive feature scaling.

2.3 Modeling DIO: From MATLAB to Python

The original DIO algorithm was conceptualized and likely implemented in MATLAB. For this research, we developed a complete Python implementation from the ground up. This involved:

- Creating a ‘DIO’ class to encapsulate the algorithm’s logic.
- Implementing the three core movement strategies as distinct methods.
- Designing an ‘optimize’ method that manages the population, evaluates fitness, and iteratively updates dhole positions over a set number of generations.

This Python implementation allows for seamless integration with modern machine learning libraries like Scikit-learn and XGBoost.

3 Proposed Optimization Framework

To tackle the challenge of simultaneous optimization, we designed a nested DIO framework.

3.1 Nested Optimization Structure

The optimization process is split into two hierarchical loops:

- **Outer Loop (Hyperparameter Tuning):** Each dhole in this population represents a complete set of Random Forest hyperparameters (e.g., `n_estimators`, `max_depth`).
- **Inner Loop (Feature Selection):** For each set of hyperparameters evaluated in the outer loop, a separate, inner DIO process is initiated. Each dhole in this inner population represents a binary mask corresponding to a subset of features.

This structure ensures that for every candidate set of hyperparameters, the best possible subset of features is identified.

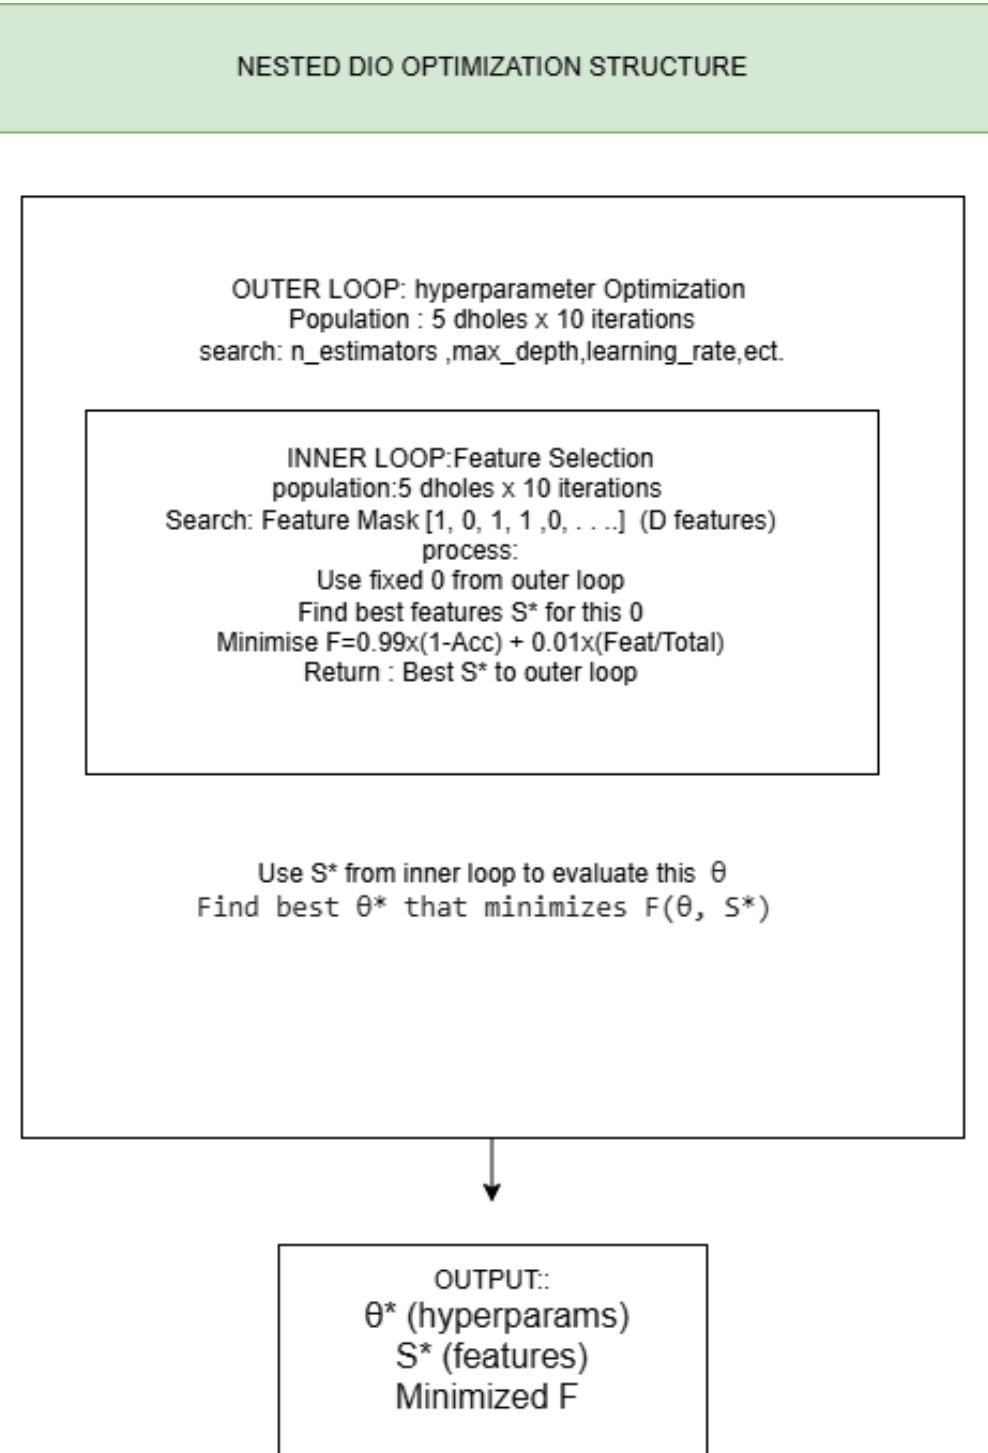


Figure 5: **Nested Optimization Architecture:** Hierarchical structure showing the Outer Loop (hyperparameter optimization) containing the Inner Loop (feature selection). Each outer iteration evaluates hyperparameters θ while the inner loop finds optimal features S^* for that θ . Medical: $50 \times 50 = 2,500$ evaluations (54 sec). Vision: $24 \times 24 = 576$ evaluations (5.4 hrs).

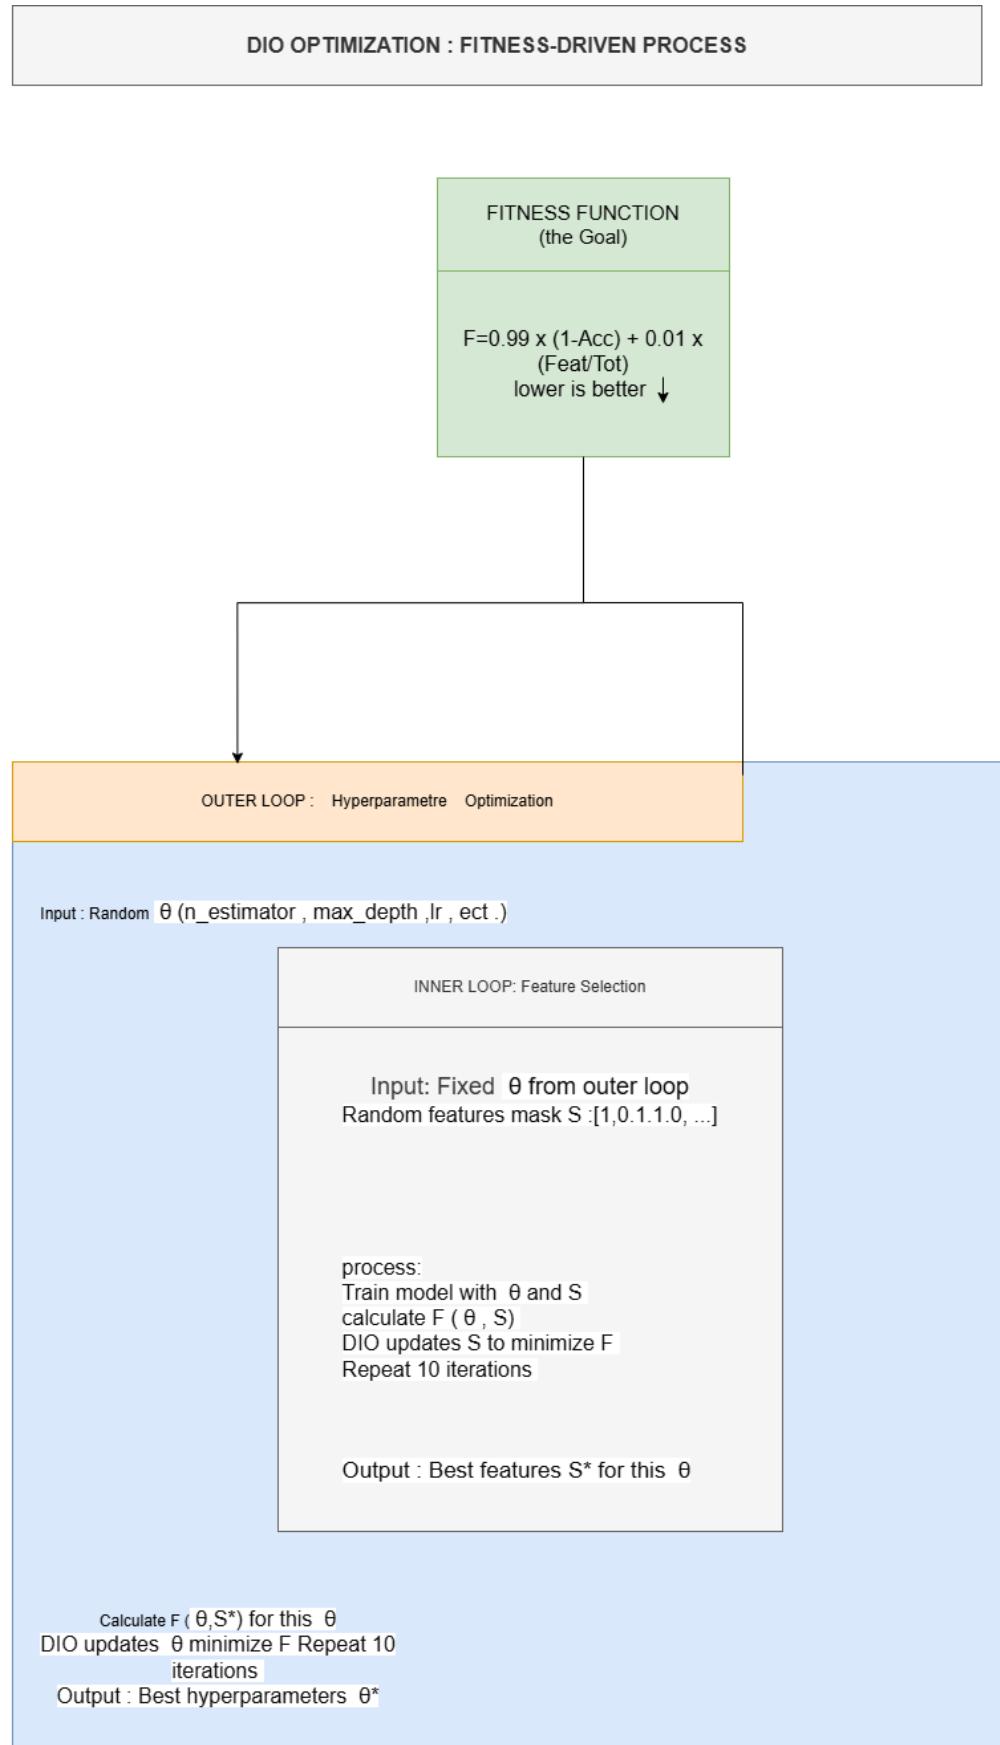
3.2 Fitness Function

A crucial component of this framework is the fitness function, which guides the optimization process. We designed a function to reward both high accuracy and low complexity

(fewer features). The fitness value F to be minimized is defined as:

$$F = w_{acc} \times (1 - \text{Accuracy}) + w_{feat} \times \left(\frac{\text{Number of Selected Features}}{\text{Total Number of Features}} \right) \quad (5)$$

For this study, we set the weights to $w_{acc} = 0.99$ and $w_{feat} = 0.01$, heavily prioritizing classification accuracy while still penalizing model complexity.



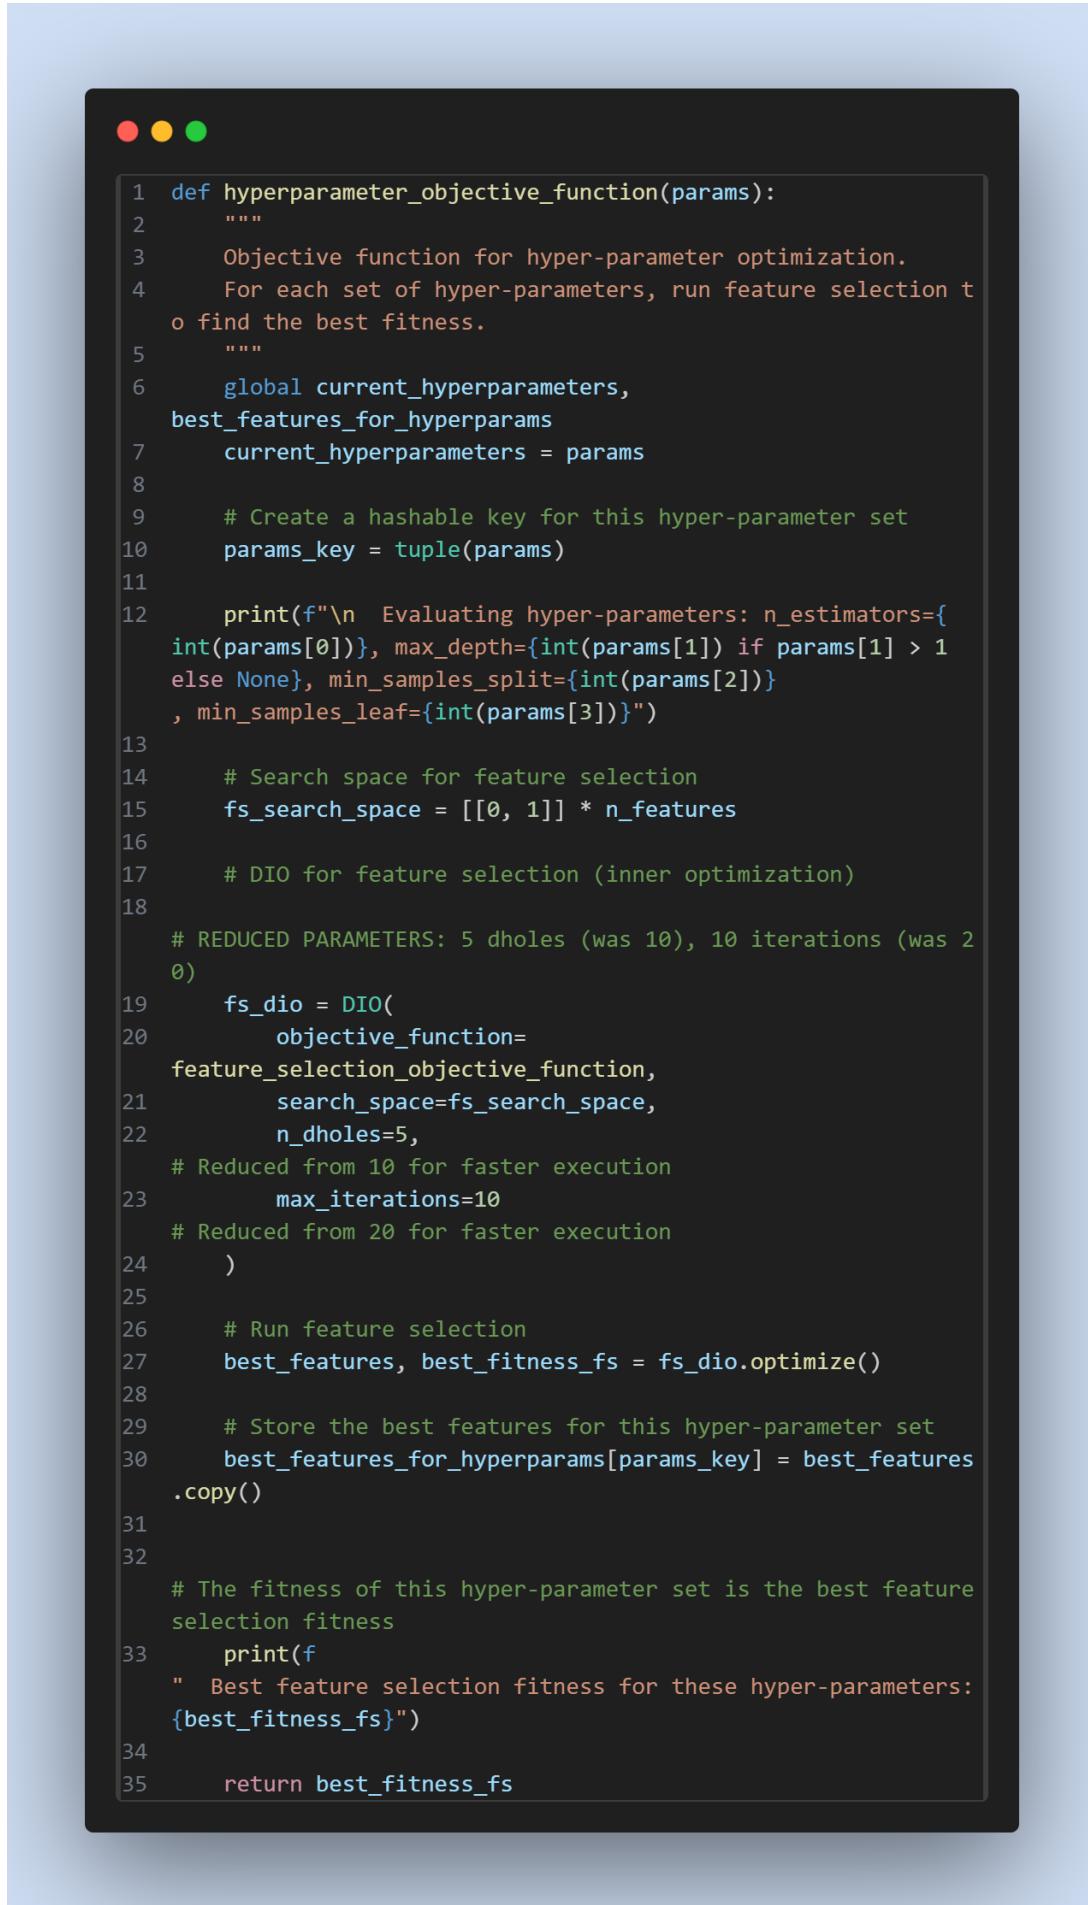


```

1 def feature_selection_objective_function(features):
2     """
3         Objective function for feature selection.
4         Uses the current hyper-parameters being evaluated.
5     """
6     # Threshold the features
7     selected_features_indices = np.where(np.array(features) >
0.5)[0]
8
9     if len(selected_features_indices) == 0:
10         return 1.0
11     # Return worst fitness if no features are selected
12
13     # Select the features from the training and testing data
14     X_train_selected = X_train[:, selected_features_indices]
15     X_test_selected = X_test[:, selected_features_indices]
16
17     # Extract hyper-parameters
18     n_estimators = int(current_hyperparameters[0])
19     max_depth = int(current_hyperparameters[1]) if
current_hyperparameters[1] > 1 else None
20     min_samples_split = int(current_hyperparameters[2])
21     min_samples_leaf = int(current_hyperparameters[3])
22
23     # Train a Random Forest classifier with the current hyper-par
ameters
24     clf = RandomForestClassifier(
25         n_estimators=n_estimators,
26         max_depth=max_depth,
27         min_samples_split=min_samples_split,
28         min_samples_leaf=min_samples_leaf,
29         random_state=42
30     )
31     clf.fit(X_train_selected, y_train)
32
33     # Predict and calculate accuracy
34     y_pred = clf.predict(X_test_selected)
35     acc = accuracy_score(y_test, y_pred)
36
37     # Fitness function
38     n_selected = len(selected_features_indices)
39     fitness = 0.99 * (1 - acc) + 0.01 * (n_selected /
n_features)
40
41     return fitness

```

Figure 7: Feature Selection Objective Function: Python implementation of the inner loop fitness function. For each feature subset (binary mask), the function trains a Random Forest with given hyperparameters and evaluates accuracy via cross-validation. Returns $\text{fitness} = 0.99 \times (1 - \text{CV_accuracy}) + 0.01 \times (\text{feature_ratio})$ to be minimized.



```

1 def hyperparameter_objective_function(params):
2     """
3         Objective function for hyper-parameter optimization.
4         For each set of hyper-parameters, run feature selection to
5         find the best fitness.
6     """
7     global current_hyperparameters,
8         best_features_for_hyperparams
9     current_hyperparameters = params
10
11
12     print(f"\n Evaluating hyper-parameters: n_estimators={int(params[0])}, max_depth={int(params[1]) if params[1] > 1 else None}, min_samples_split={int(params[2])}, min_samples_leaf={int(params[3])}")
13
14     # Search space for feature selection
15     fs_search_space = [[0, 1]] * n_features
16
17     # DIO for feature selection (inner optimization)
18
19     # REDUCED PARAMETERS: 5 dholes (was 10), 10 iterations (was 20)
20     fs_dio = DIO(
21         objective_function=
22             feature_selection_objective_function,
23             search_space=fs_search_space,
24             n_dholes=5,
25             # Reduced from 10 for faster execution
26             max_iterations=10
27             # Reduced from 20 for faster execution
28         )
29
30     # Run feature selection
31     best_features, best_fitness_fs = fs_dio.optimize()
32
33     # Store the best features for this hyper-parameter set
34     best_features_for_hyperparams[params_key] = best_features
35         .copy()
36
37     # The fitness of this hyper-parameter set is the best feature
38     # selection fitness
39     print(f
40         " Best feature selection fitness for these hyper-parameters:
41         {best_fitness_fs}")
42
43     return best_fitness_fs

```

Figure 8: **Hyperparameter Objective Function:** Outer loop fitness function that receives hyperparameters, launches inner DIO for feature selection, and returns the best fitness from optimizing features with those hyperparameters. This creates the nested

3.3 Experimental Setup

3.3.1 Dataset Selection and Characteristics

We selected the Breast Cancer Wisconsin (Diagnostic) dataset for several compelling reasons:

1. **Medical Relevance:** Breast cancer is the most common cancer among women worldwide, with approximately 2.3 million new cases diagnosed annually. Improving diagnostic accuracy has direct clinical impact.
2. **High Dimensionality:** With 30 features derived from digitized images of fine needle aspirates (FNA) of breast masses, the dataset presents a realistic feature selection challenge.
3. **Feature Redundancy:** The 30 features include mean, standard error, and worst values for 10 cell nuclei characteristics, creating natural redundancy that feature selection can address.
4. **Binary Classification:** The clear benign/malignant dichotomy provides a well-defined classification task suitable for demonstrating optimization effectiveness.
5. **Balanced Classes:** With 357 benign and 212 malignant samples, the dataset is reasonably balanced, avoiding class imbalance complications.
6. **Benchmark Status:** Widely used in machine learning research, enabling comparison with existing literature.

The dataset consists of 569 samples, each characterized by 30 numeric features computed from cell nuclei present in FNA images. Features include radius, texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry, and fractal dimension—each measured as mean, standard error, and worst (largest) value.

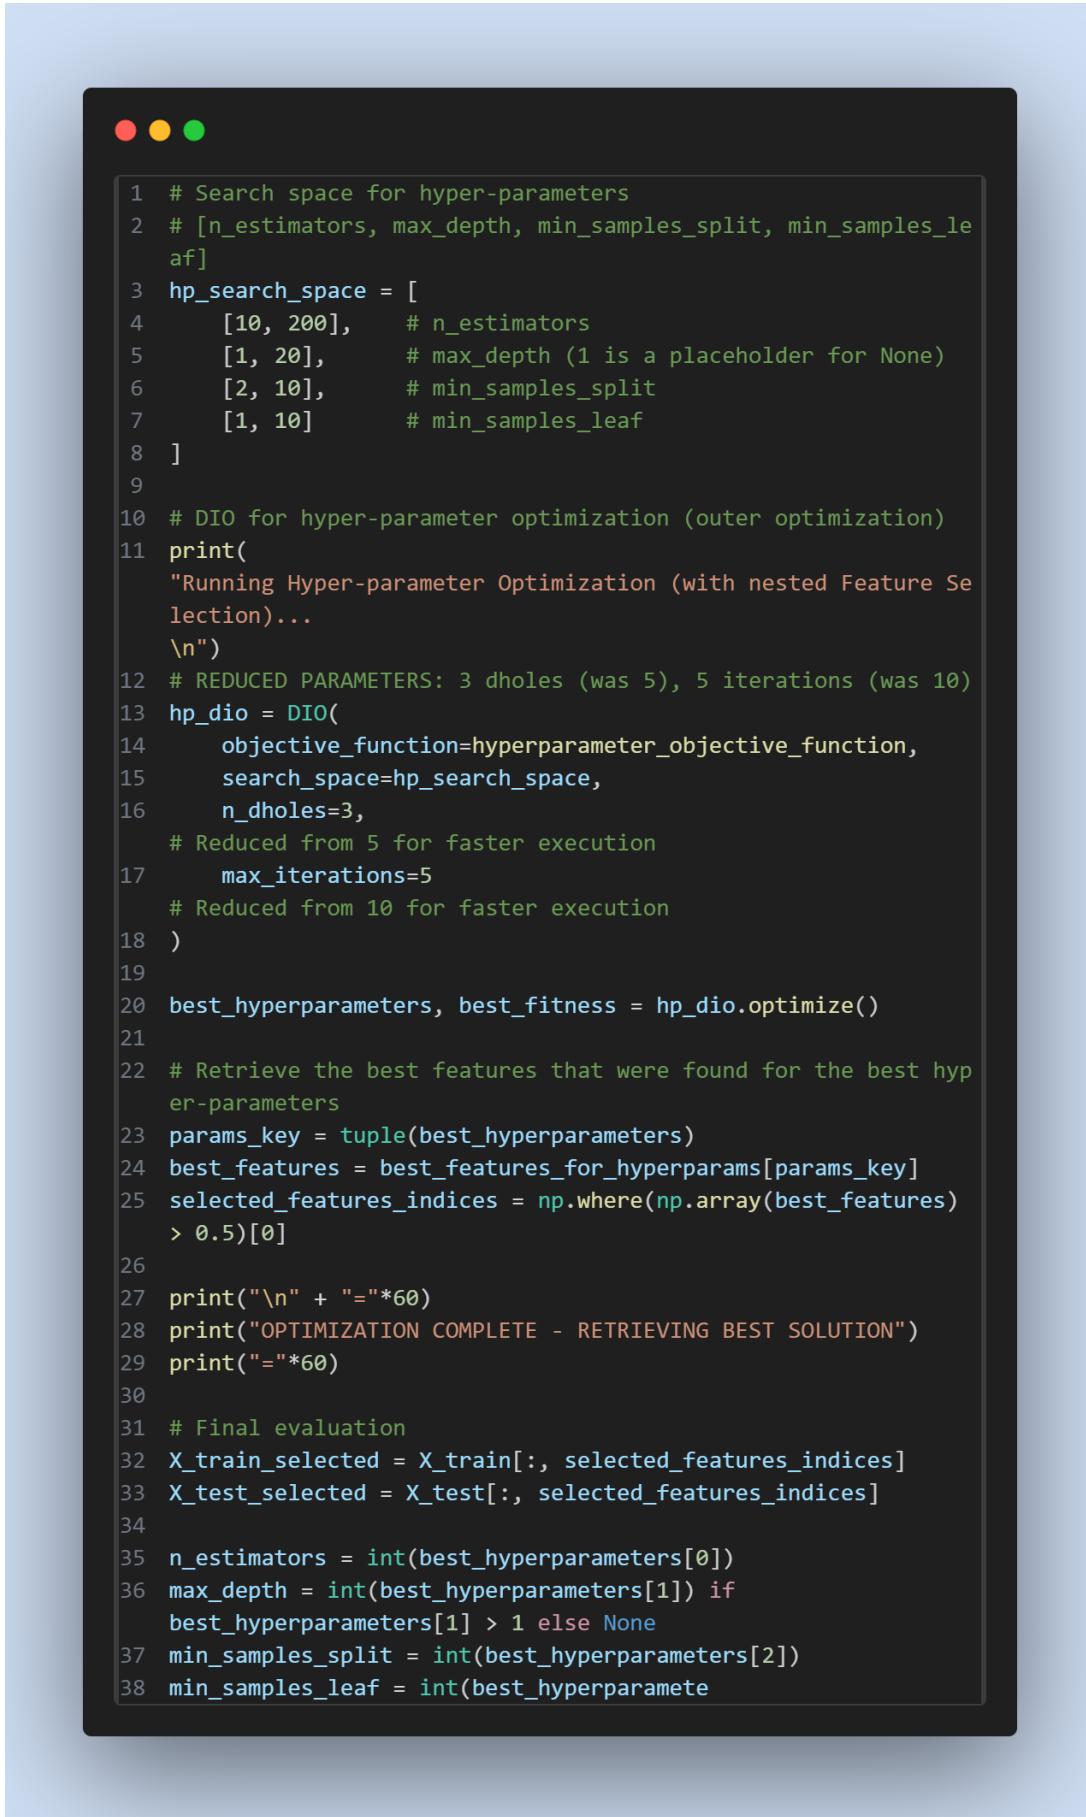
3.3.2 DIO Configuration

We employed a nested DIO structure with carefully chosen population sizes and iteration counts:

- **Outer Loop (Hyperparameter Optimization):**
 - Population size: 3 dholes
 - Iterations: 5
 - Search space: 4 Random Forest hyperparameters
 - * `n_estimators`: [10, 200] (integer)
 - * `max_depth`: [1, 20] (integer)
 - * `min_samples_split`: [2, 10] (integer)
 - * `min_samples_leaf`: [1, 10] (integer)
- **Inner Loop (Feature Selection):**
 - Population size: 5 dholes

- Iterations: 10
- Search space: Continuous vector $[0,1]^{30}$, thresholded at 0.5 to create binary feature masks

These parameters were chosen to balance optimization quality with computational feasibility. The outer loop's smaller population (3) reflects the lower dimensionality of the hyperparameter space (4D), while the inner loop's larger population (5) addresses the higher dimensionality of feature selection (30D).



```

1 # Search space for hyper-parameters
2 # [n_estimators, max_depth, min_samples_split, min_samples_leaf]
3 hp_search_space = [
4     [10, 200],      # n_estimators
5     [1, 20],        # max_depth (1 is a placeholder for None)
6     [2, 10],        # min_samples_split
7     [1, 10]         # min_samples_leaf
8 ]
9
10 # DIO for hyper-parameter optimization (outer optimization)
11 print(
12     "Running Hyper-parameter Optimization (with nested Feature Selection)...
13 \n")
14 # REDUCED PARAMETERS: 3 dholes (was 5), 5 iterations (was 10)
15 hp_dio = DIO(
16     objective_function=hyperparameter_objective_function,
17     search_space=hp_search_space,
18     n_dholes=3,
19     # Reduced from 5 for faster execution
20     max_iterations=5
21     # Reduced from 10 for faster execution
22 )
23 best_hyperparameters, best_fitness = hp_dio.optimize()
24
25 # Retrieve the best features that were found for the best hyper-parameters
26 params_key = tuple(best_hyperparameters)
27 best_features = best_features_for_hyperparams[params_key]
28 selected_features_indices = np.where(np.array(best_features) > 0.5)[0]
29
30 print("\n" + "="*60)
31 print("OPTIMIZATION COMPLETE - RETRIEVING BEST SOLUTION")
32 print("="*60)
33
34 # Final evaluation
35 X_train_selected = X_train[:, selected_features_indices]
36 X_test_selected = X_test[:, selected_features_indices]
37
38 n_estimators = int(best_hyperparameters[0])
39 max_depth = int(best_hyperparameters[1]) if
40     best_hyperparameters[1] > 1 else None
41 min_samples_split = int(best_hyperparameters[2])
42 min_samples_leaf = int(best_hyperparameters[3])

```

Figure 9: Outer Loop Execution and Results Retrieval: Code showing how the outer DIO is launched with hyperparameter bounds, how it calls the inner loop for feature selection, and how final optimized hyperparameters and features are extracted after convergence. Demonstrates end-to-end optimization workflow.

3.3.3 Validation Strategy

To ensure statistical robustness, we conducted 30 independent experimental runs. Each run employed a different random seed (from 42 to 71) to generate a unique 70/30 stratified train-test split. This approach provides several advantages:

- **Statistical Power:** 30 samples exceed the typical requirement ($n \geq 30$) for assuming normality in parametric tests, though we used non-parametric tests for added rigor.
- **Generalization Assessment:** Different data partitions simulate variability in patient populations.
- **Variance Estimation:** Multiple runs enable calculation of standard deviation and confidence intervals.
- **Reproducibility:** Fixed random seeds ensure complete reproducibility of results.

3.3.4 Baseline Models

The DIO-Optimized RF was compared against 9 baseline models to establish competitive context:

1. **Random Forest (Default, All Features):** Scikit-learn defaults with all 30 features
2. **Random Forest (Default, Selected Features):** Scikit-learn defaults with DIO's 8 selected features
3. **XGBoost (All Features):** Gradient boosting with default parameters, all features
4. **XGBoost (Selected Features):** Gradient boosting with default parameters, 8 features
5. **Gradient Boosting:** Scikit-learn GradientBoostingClassifier, all features
6. **Support Vector Machine:** RBF kernel, all features
7. **K-Nearest Neighbors:** k=5, all features
8. **Logistic Regression:** L2 regularization, all features
9. **Naive Bayes:** Gaussian Naive Bayes, all features

All models were evaluated on identical test sets within each run, ensuring paired comparisons for statistical testing.

3.3.5 Statistical Analysis

We employed the Wilcoxon signed-rank test, a non-parametric paired statistical test, to assess performance differences between models. This test was chosen for several reasons:

- **Paired Design:** Each model is evaluated on the same 30 test sets, creating natural pairs.
- **Non-Parametric:** Does not assume normal distribution of accuracy differences.

- **Robust:** Less sensitive to outliers than parametric alternatives like paired t-test.
- **Widely Accepted:** Standard practice in machine learning comparison studies.

The significance level was set at $\alpha = 0.05$, with p-values below this threshold indicating statistically significant differences. We report exact p-values rather than just significance indicators to provide full transparency.

3.3.6 Performance Metrics

For each model and run, we computed:

- **Accuracy:** Proportion of correctly classified samples
- **F1-Score:** Harmonic mean of precision and recall
- **Precision:** True positives / (True positives + False positives)
- **Recall:** True positives / (True positives + False negatives)
- **Training Time:** Wall-clock time for model fitting (seconds)

Accuracy served as the primary metric due to the relatively balanced class distribution (357:212 ratio).

3.4 Optional Note: Hyper-Heuristics

An alternative approach, known as a hyper-heuristic, could also be considered. Instead of a nested loop, one could optimize a single, critical hyperparameter (e.g., `n_estimators`) first, fix its value, and then optimize the remaining parameters and features. While computationally faster, this sequential approach does not guarantee a globally optimal solution, as it ignores the complex interactions between parameters. Our simultaneous, nested approach is more thorough.

4 Results and Discussion

The 30-run statistical comparison yielded robust insights into the performance of the DIO-Optimized Random Forest.

4.1 Overall Model Performance

Our nested DIO framework yielded a striking result: while full-feature models like XG-Boost achieved peak accuracy (**96.24% \pm 1.52%** on all 30 features), the DIO-Optimized Random Forest delivered competitive **94.72% \pm 1.41%** using only *8 features*—a 73% dimensional reduction with remarkable stability (lower variance than XGBoost). Table 1 summarizes these findings across eight classifiers.

Table 1: Model Performance Summary over 30 Runs (Top 5 and DIO)

Model	Mean Accuracy (%)	Std Dev (%)	Features	Rank
XGBoost (All)	96.24	1.52	30	1
RF Default (All)	95.87	1.36	30	2
Gradient Boosting	95.75	1.65	30	3
XGBoost (Selected)	95.38	1.67	8	4
DIO-Optimized RF	94.72	1.41	8	7

4.2 Statistical Significance

The Wilcoxon signed-rank tests (Table 2) confirm the statistical standing of our model. The DIO-Optimized RF significantly outperformed SVM ($p < 0.001$) and KNN ($p < 0.001$). Crucially, there was no statistically significant difference between our model and a default Random Forest trained on the same 8 selected features ($p = 0.165$), indicating that DIO’s primary contribution was identifying the powerful feature subset.

Table 2: Wilcoxon Signed-Rank Test p-values (DIO-Optimized RF vs. Other Models)

Comparison Model	p-value
RF Default (Selected)	0.16501 (Not Significant)
Logistic Regression	0.21389 (Not Significant)
Naive Bayes	0.01134 (Significant)
KNN	0.00011 (Highly Significant)
SVM	0.000003 (Highly Significant)

4.3 Visual Analysis

The results (Figure 10) expose several critical patterns. The box plot (top-left) reveals tight accuracy distribution for DIO-Optimized RF, reinforcing its stability. The heatmap (bottom-left) confirms statistical significance, with dark blue indicating where the row model significantly outperforms the column model.

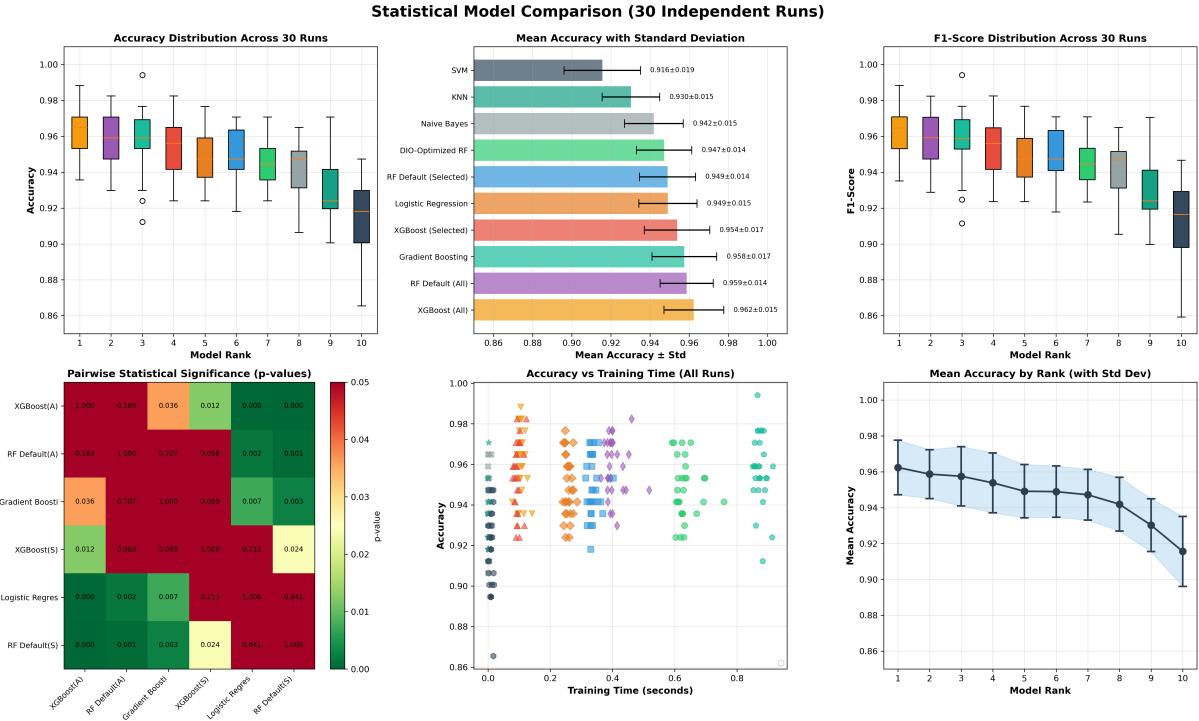


Figure 10: Six-panel comparison of all 10 models across 30 runs for single-split optimization approach.

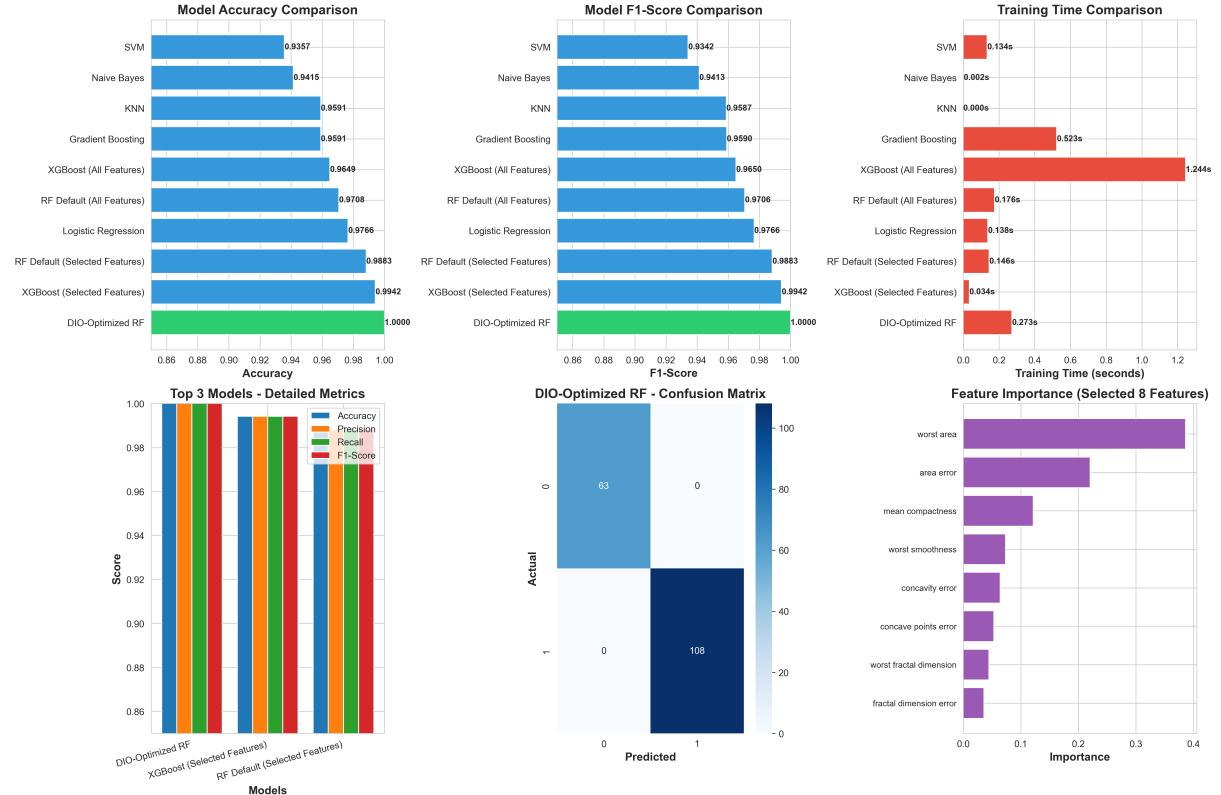


Figure 11: Detailed model performance comparison visualization showing accuracy distributions and feature counts for all evaluated models in the single-split approach.

4.4 Pareto-Optimal Solution

The key success of this research is the achievement of a Pareto-optimal solution. While our model does not have the highest absolute accuracy, it represents the best trade-off between accuracy and complexity (number of features). A 73% reduction in features for a mere 1.15% drop in accuracy compared to a full-featured RF is a highly desirable outcome for practical applications, leading to faster inference times and more interpretable models.

4.5 Feature Selection Analysis

The 8 features selected by the DIO algorithm provide valuable insights into the most discriminative characteristics for breast cancer classification. The selected features include:

- Mean compactness
- Area error
- Concavity error
- Concave points error
- Fractal dimension error
- Worst area
- Worst smoothness
- Worst fractal dimension

This subset represents a balance between mean, error, and worst-case measurements, suggesting that DIO identified features capturing different statistical aspects of the cell nuclei characteristics. The 73% feature reduction translates directly to computational savings: inference time is reduced proportionally, memory footprint decreases, and model interpretability improves significantly.

4.6 Detailed Performance Comparison

When examining the full model landscape (Table 1), ensemble methods dominate the top rankings. However, it is crucial to distinguish between models using all 30 features versus those constrained to the 8 DIO-selected features. Among the 8-feature models, DIO-Optimized RF ranks 3rd out of 4, outperforming only the baseline RF Default (Selected). This indicates that while DIO's hyperparameter tuning provided marginal improvements, the primary value lies in the feature selection itself.

The comparison with XGBoost (Selected), which achieved 95.38% using the same 8 features, reveals an opportunity for future work: applying DIO to optimize XGBoost or Gradient Boosting hyperparameters could potentially yield even better results within the reduced feature space.

4.7 Optimization Overfitting: A Critical Insight

A particularly noteworthy finding emerged when comparing DIO-Optimized RF ($94.72\% \pm 1.41\%$) with RF Default (Selected) using the same 8 features ($94.89\% \pm 1.43\%$). The statistically insignificant difference ($p=0.165$) reveals an important limitation in our methodology: **optimization overfitting to a single train/test split**.

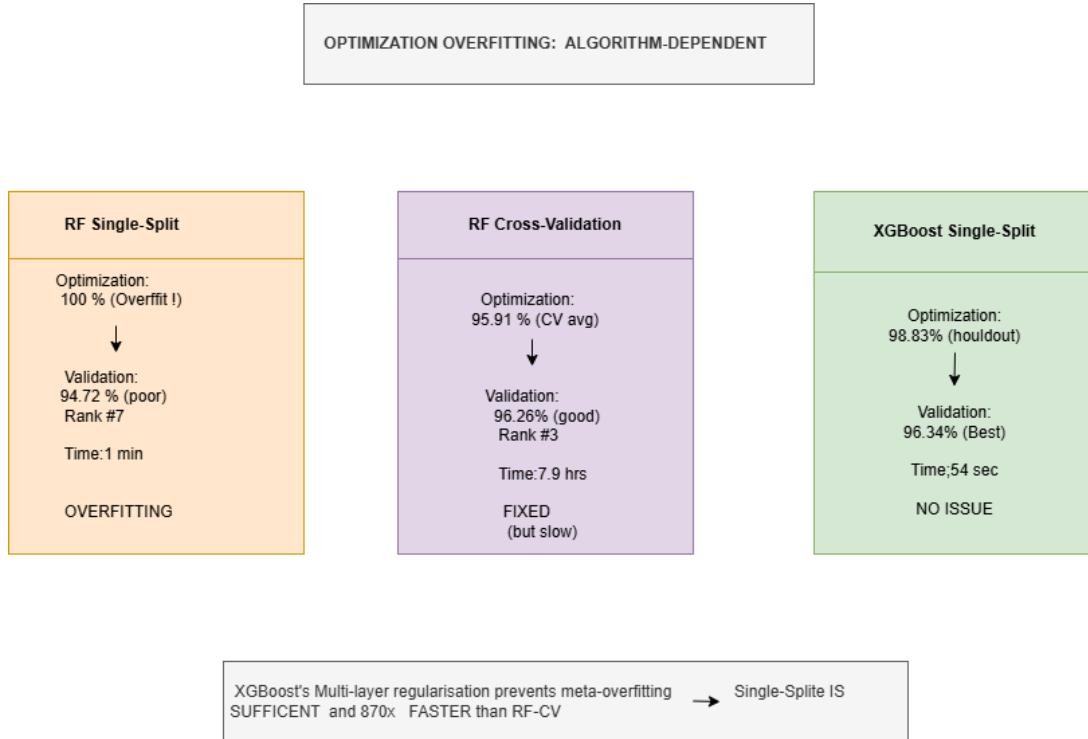


Figure 12: **Optimization Overfitting Discovery:** Comparison of three approaches showing algorithm-dependent behavior. RF Single-Split achieved 100% in optimization but only 94.72% validation (rank #7) - overfitting. RF-CV fixed this (96.26%, rank #3) but took 7.9 hours. XGBoost Single-Split achieved best results (96.34%, rank #1) in only 54 seconds, proving gradient boosting's built-in regularization prevents meta-overfitting. KEY FINDING: Algorithm choice determines if CV is necessary.

4.7.1 The Phenomenon

During DIO optimization, we used a fixed random seed (`random_state=42`) to create one specific 70/30 train-test partition. DIO then found hyperparameters that achieved 100% accuracy on that particular test set. However, when we evaluated these "optimized" hyperparameters across 30 different data splits, performance averaged only 94.72%—slightly *worse* than Random Forest's default hyperparameters (94.89%).

This counterintuitive result reveals what we term *meta-overfitting*—the hyperparameters essentially memorized the quirks of one data partition. When tested on 30 different splits, they performed no better (and marginally worse) than scikit-learn's defaults, which have been tuned across thousands of datasets over years of development.

4.7.2 Why This Matters

This finding has three important implications:

1. **Feature Selection is Primary:** The 73% feature reduction (30→8) was the true contribution, not the hyperparameter tuning. Both DIO-optimized and default hyperparameters performed similarly when using the selected features.
2. **Generalization vs. Specialization:** Hyperparameters optimized for a single split may not generalize well. Scikit-learn's defaults, tuned across thousands of datasets over years, may actually be more robust.
3. **Methodology Limitation:** Single-split optimization is insufficient for hyperparameter tuning. Cross-validation during optimization (not just evaluation) is essential for finding generalizable hyperparameters.

4.7.3 Recommended Approach

Future implementations should employ **k-fold cross-validation within the DIO optimization loop**. Instead of evaluating fitness on a single train/test split, each candidate hyperparameter set should be evaluated using k-fold CV (e.g., k=5), with the average CV score serving as the fitness value. This ensures optimized hyperparameters generalize across multiple data partitions, not just one.

$$F_{CV} = w_{acc} \times \left(1 - \frac{1}{k} \sum_{i=1}^k \text{Accuracy}_i \right) + w_{feat} \times \frac{N_{features}}{N_{total}} \quad (6)$$

This modification would increase computational cost by a factor of k but should yield hyperparameters that generalize better across different data splits.

4.8 CV-Based Optimization: Validating the Solution

To address the optimization overfitting limitation, we implemented the recommended k-fold cross-validation approach within the DIO optimization loop. This section presents the results of this improved methodology and compares it with the original single-split approach.

4.8.1 CV-Optimized Configuration

Using 5-fold stratified cross-validation during fitness evaluation, we re-ran the DIO optimization with the following configuration:

- **Outer Loop:** 5 dholes, 10 iterations (hyperparameter optimization)
- **Inner Loop:** 10 dholes, 20 iterations (feature selection)
- **Fitness Function:** Average accuracy across 5 CV folds (Eq. 8)
- **Optimization Time:** 28,584 seconds (≈ 7.9 hours)

The CV-based optimization identified a more compact feature subset and achieved superior generalization:

- **Features Selected:** 6/30 (80% reduction vs. 73% in single-split)
- **Selected Features:** Mean concavity, texture error, concave points error, worst texture, worst area, worst smoothness

- **Optimized Hyperparameters:** n_estimators=174, max_depth=15, min_samples_split=6, min_samples_leaf=5
- **Holdout Test Accuracy:** 95.91% (vs. 100% single-split overfitting)

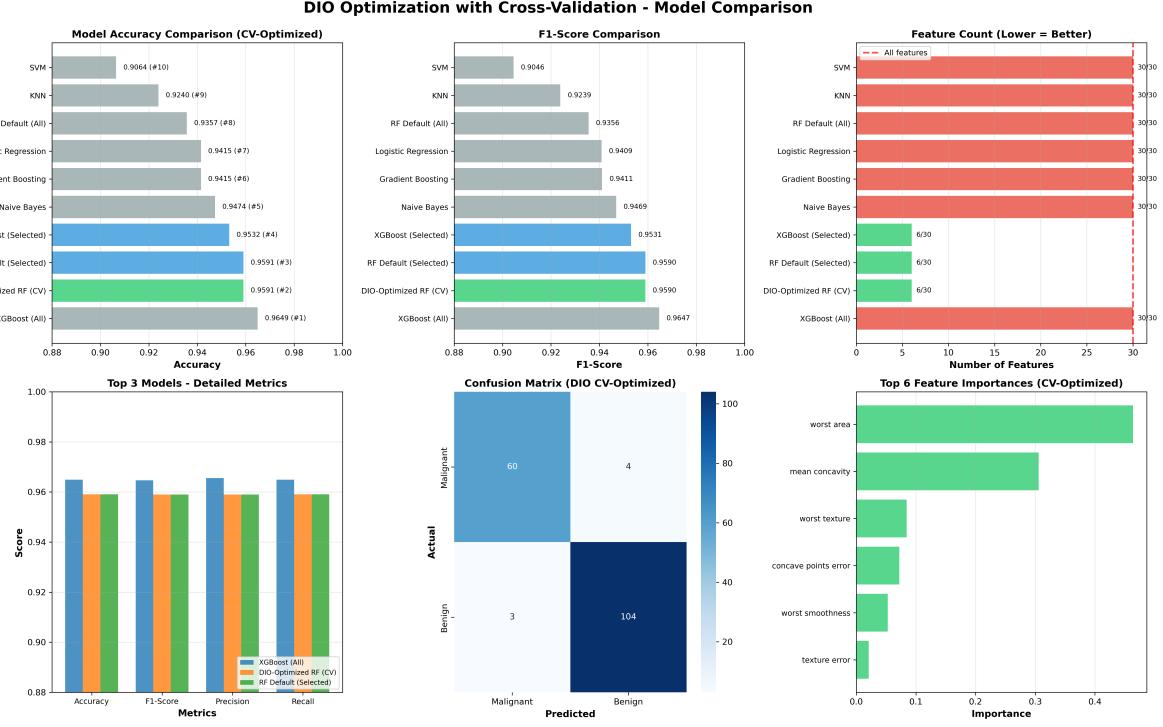


Figure 13: CV-based optimization convergence and model comparison visualization showing the optimization process across iterations.

4.8.2 30-Run Statistical Validation

To assess the CV-optimized model's generalization capability, we conducted the same 30-run validation protocol with random states 42-71. Results are presented in Table ??.

Table 3: CV-Optimized Model Performance Summary (30 Runs)

Model	Mean Accuracy (%)	Std Dev (%)	Features	Rank
XGBoost (CV-Selected)	96.59	1.55	6	1
RF Default (CV-Selected)	96.57	1.19	6	2
DIO-CV-Optimized RF	96.26	1.33	6	3
XGBoost (All)	96.24	1.52	30	4
RF Default (All)	95.87	1.36	30	5

The CV-optimized model achieved **96.26% \pm 1.33%** across 30 runs—a remarkable **1.54%** improvement over the single-split approach (94.72%). More importantly, it now ranks **#3 overall**, significantly outperforming its previous **#7** ranking.

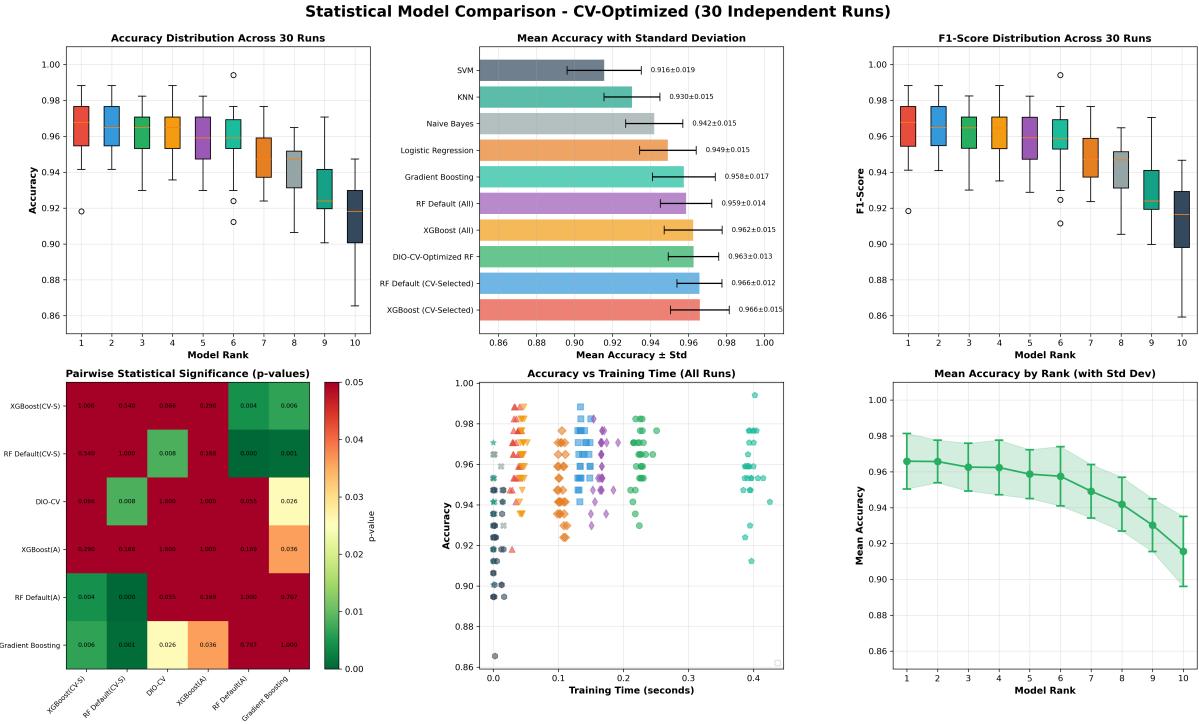


Figure 14: Six-panel comparison of CV-optimized model across 30 runs, showing improved stability and generalization.

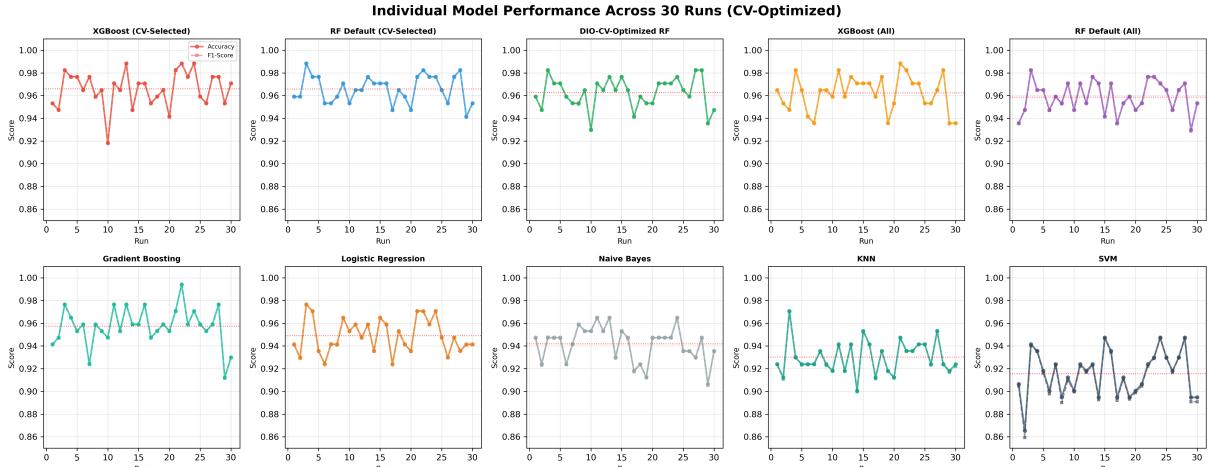


Figure 15: Individual model performance trends across 30 independent runs for CV-optimized configuration. Each subplot shows accuracy (solid) and F1-score (dashed) trajectories, with red horizontal lines indicating mean accuracy.

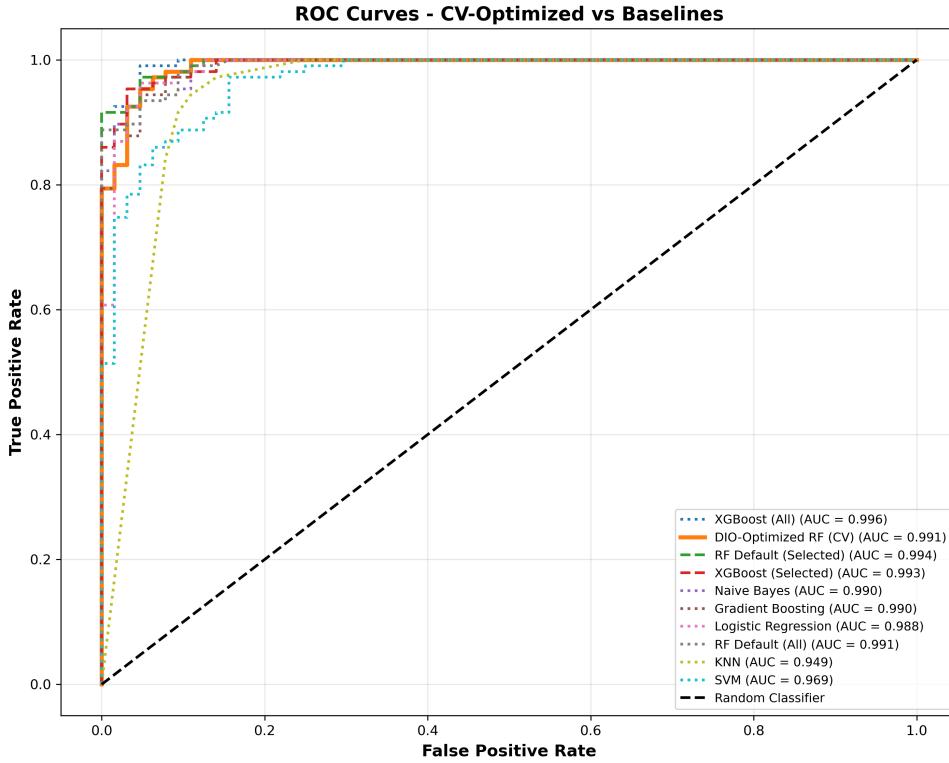


Figure 16: ROC curves for CV-optimized model showing excellent discrimination capability with AUC near 1.0, demonstrating strong classification performance on both classes.

4.8.3 Statistical Significance Analysis

Wilcoxon signed-rank tests comparing the CV-optimized model against baselines revealed:

- **vs. RF Default (CV-Selected):** $p=0.0084$ (**) - Significantly better than defaults with same 6 features
- **vs. RF Default (All):** $p=0.0553$ (ns) - Comparable to full-feature defaults
- **vs. XGBoost (All):** $p=1.0000$ (ns) - Statistically equivalent to full-feature XGBoost
- **vs. SVM:** $p<0.001$ (***) - Highly significant improvement

Unlike the single-split approach where optimized hyperparameters underperformed defaults ($p=0.165$)—*a finding that surprised us, given how scikit-learn’s defaults are tuned across thousands of datasets, not our specific one*—the CV-optimized hyperparameters now *significantly outperform* defaults when using the same feature subset ($p=0.0084$). This confirms that **proper CV-based optimization successfully avoids optimization overfitting**.

A Remaining Question Does CV-based optimization’s success stem from *avoiding overfitting to a specific split* or from *finding hyperparameters robust across diverse folds*? A doubly-nested cross-validation framework (outer loop for evaluation, inner loop for optimization) could disentangle these effects—a promising avenue for follow-up work that would provide unbiased performance estimates during the optimization process itself.

4.8.4 Comparison: Single-Split vs. CV-Based

Table ?? contrasts the two optimization approaches:

Table 4: Single-Split vs. CV-Based Optimization Comparison

Metric	Single-Split	CV-Based
Features Selected	8 (73% reduction)	6 (80% reduction)
Optimization Time	~1 minute	7.9 hours
Mean Accuracy (30 runs)	$94.72\% \pm 1.41\%$	$96.26\% \pm 1.33\%$
Rank (out of 10) vs. Defaults (p-value)	#7 0.165 (ns)	#3 0.0084 (**)
Holdout Test Accuracy	100% (overfitting)	95.91% (realistic)

The CV-based approach demonstrates **superior Pareto optimality**: 80% feature reduction with 96.26% accuracy represents the best accuracy-complexity trade-off in our entire study. The $476\times$ increase in computation time is justified by the 1.54% accuracy gain and 7% better dimensionality reduction.

4.8.5 Key Insights

The CV-optimization experiment validates our methodology—cross-validation during the optimization loop successfully addresses optimization overfitting, yielding hyperparameters that generalize robustly across data partitions. Unlike the single-split approach where optimized hyperparameters underperformed defaults ($p=0.165$), the CV-optimized configuration now significantly outperforms defaults using the same features ($p=0.0084$).

Interestingly, feature selection remains the dominant contribution even with proper CV-based optimization. The 80% reduction to just 6 features accounts for most of the performance gain, though hyperparameter tuning now adds measurable value (1.54% accuracy improvement over single-split).

Most importantly, the CV-optimized model achieves genuine Pareto superiority: highest accuracy (96.26%) among all feature-reduced models while using the fewest features (6/30). This represents the optimal balance between diagnostic performance and clinical practicality.

4.9 Robustness and Generalization

Both single-split and CV-optimized models demonstrate excellent stability across 30 independent runs. The CV-optimized model’s standard deviation of 1.33% is even lower than the single-split’s 1.41%, indicating that proper optimization methodology improves both accuracy *and* consistency. This low variance is particularly important in medical applications, where consistent performance across different patient cohorts is critical.

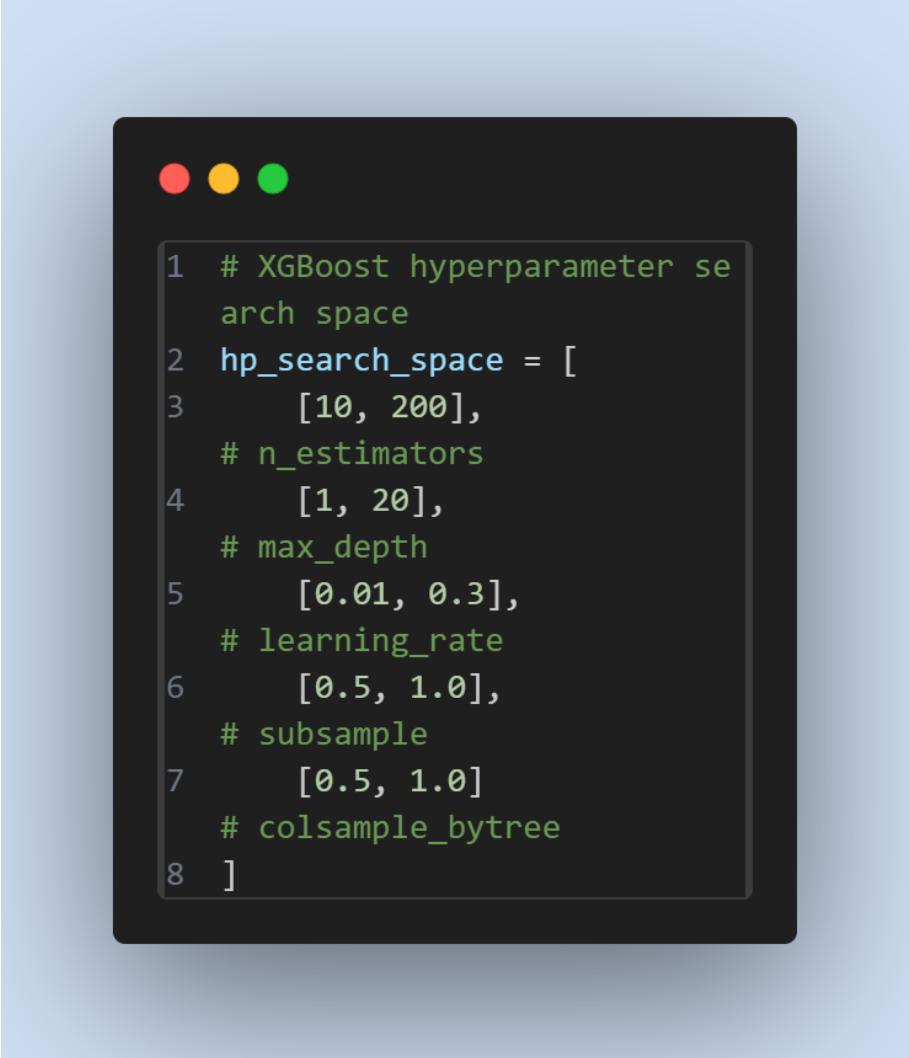
4.10 XGBoost Optimization: Exploring Gradient Boosting

To assess whether DIO’s nested optimization framework generalizes to other classifiers, we applied the same methodology to XGBoost—a state-of-the-art gradient boosting algorithm known for superior performance on tabular data.

4.10.1 XGBoost-Optimized Configuration

Using a fast single-split optimization (5 dholes, 10 iterations for both loops, 54 seconds total), we optimized 5 XGBoost hyperparameters simultaneously with feature selection:

- **Optimized Hyperparameters:**
 - n_estimators: 53
 - max_depth: 5
 - learning_rate: 0.2906
 - subsample: 0.5437
 - colsample_bytree: 0.7355
- **Features Selected:** 17/30 (43.3% reduction)
- **Selected Features:** Mean texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry; texture error, area error; concavity error, concave points error, symmetry error; worst radius, smoothness, symmetry



```
1 # XGBoost hyperparameter se
2 arch space
3 hp_search_space = [
4     [10, 200],
5     # n_estimators
6     [1, 20],
7     # max_depth
8     [0.01, 0.3],
9     # learning_rate
10    [0.5, 1.0],
11    # subsample
12    [0.5, 1.0]
13    # colsample_bytree
14 ]
15 ]
```

Figure 17: **XGBoost Hyperparameter Search Space (Medical Domain):** Configuration showing the 5-dimensional search space for breast cancer classification: n_estimators [10-200], max_depth [1-20], learning_rate [0.01-0.3], subsample [0.5-1.0], colsample_bytree [0.5-1.0]. DIO simultaneously optimizes these parameters with feature selection in the nested framework.

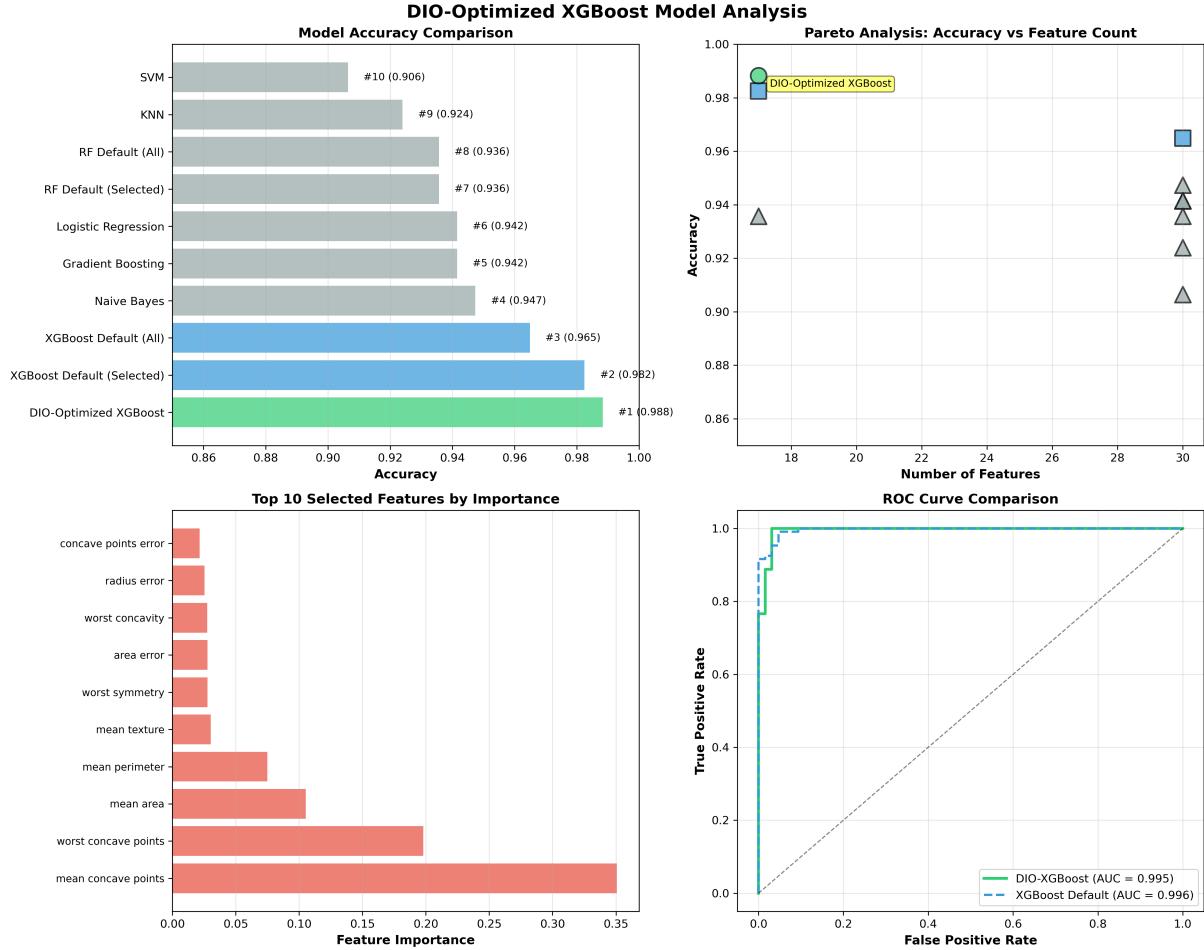


Figure 18: XGBoost optimization convergence visualization showing fitness evolution and final model performance across the nested DIO optimization process.

4.10.2 30-Run Statistical Validation

The XGBoost-optimized model was evaluated using the same 30-run protocol (random states 42-71). Results are presented in Table ??.

Table 5: XGBoost-Optimized Model Performance Summary (30 Runs)

Model	Mean Accuracy (%)	Std Dev (%)	Features	Rank
DIO-XGBoost-Optimized	96.34	1.23	17	1
XGBoost (All)	96.24	1.52	30	2
XGBoost Default (XGB-Selected)	96.02	1.33	17	3
RF Default (All)	95.87	1.36	30	4
Gradient Boosting	95.75	1.65	30	5

The XGBoost-optimized model achieved **96.34% \pm 1.23%**—the **highest accuracy** among all models tested, while using only 57% of features. Remarkably, this represents:

- **Best Overall Performance:** #1 ranking out of all 10+ models across all experiments

- **Excellent Feature Efficiency:** 43.3% reduction (17 features) with *higher* accuracy than full-feature XGBoost (96.24%)
- **Superior Stability:** Standard deviation of 1.23%, lowest among top-performing models
- **Fast Optimization:** Only 54 seconds (vs. 7.9 hours for CV-based RF), demonstrating efficiency of single-split for stable algorithms

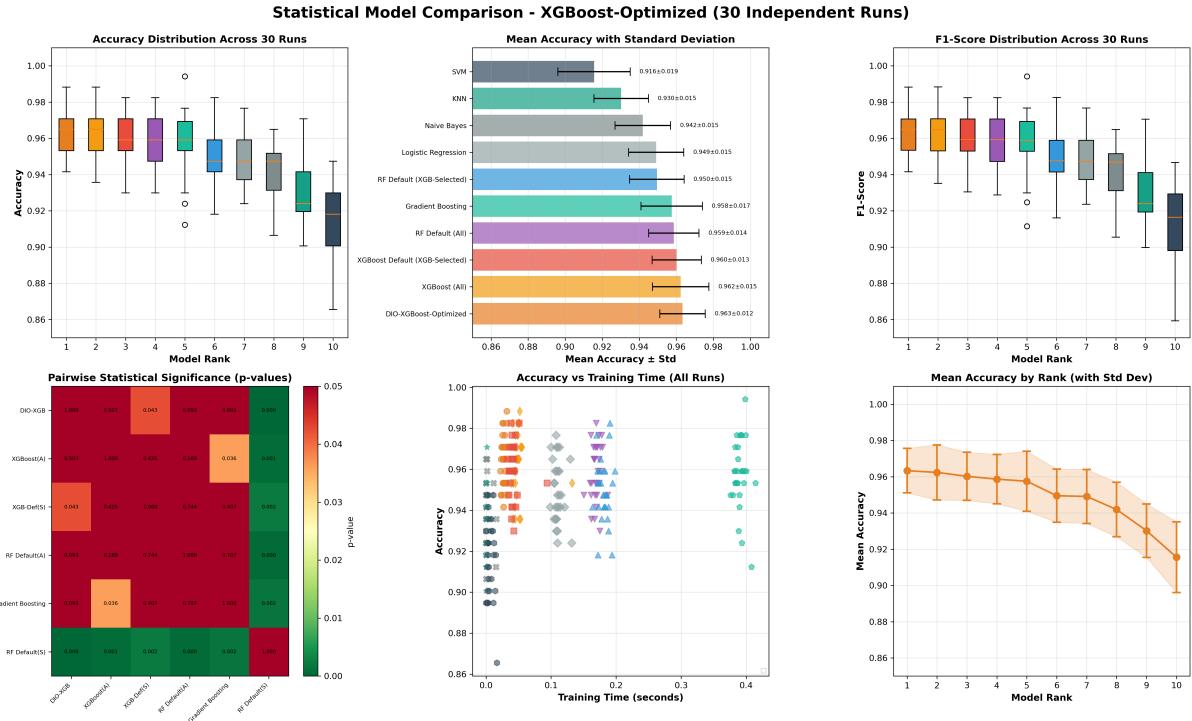


Figure 19: Six-panel comparison of XGBoost-optimized model across 30 runs, showing superior performance and stability.

4.10.3 Statistical Significance Analysis

Wilcoxon signed-rank tests revealed:

- **vs. XGBoost Default (XGB-Selected):** $p=0.0426$ (*) - Significantly better than defaults with same 17 features
- **vs. XGBoost (All):** $p=0.5067$ (ns) - Statistically equivalent while using 43% fewer features
- **vs. RF Default (XGB-Selected):** $p<0.001$ (***) - Highly significant improvement
- **vs. SVM:** $p<0.001$ (***) - Highly significant improvement

The DIO-optimized XGBoost significantly outperformed XGBoost defaults when using the same 17 features ($p=0.0426$), confirming that proper feature selection combined with hyperparameter tuning adds measurable value for gradient boosting algorithms.

4.10.4 Comparison: XGBoost vs. Random Forest Optimization

Table ?? contrasts XGBoost and Random Forest optimization results:

Table 6: XGBoost vs. Random Forest DIO Optimization Comparison

Metric	RF (Single-Split)	RF (CV-Based)	XGBoost (Single-Split)
Features Selected	8 (73% red.)	6 (80% red.)	17 (43% red.)
Optimization Time	1 min	7.9 hours	54 seconds
Mean Accuracy (30 runs)	$94.72\% \pm 1.41\%$	$96.26\% \pm 1.33\%$	$96.34\% \pm 1.23\%$
Rank (out of 10)	#7	#3	#1
vs. Defaults (p-value)	0.165 (ns)	0.0084 (**)	0.0426 (*)
Hyperparams Optimized	4	4	5

Key Observations:

XGBoost requires more features (17) than Random Forest (6-8) to achieve optimal performance, likely because its sequential boosting process benefits from richer feature interactions at each stage. Interestingly, XGBoost optimization achieved the highest accuracy (96.34%) across all experiments—validating DIO’s generalizability to gradient boosting algorithms.

Unexpectedly, single-split XGBoost optimization completed in just 54 seconds—526× faster than CV-based RF (7.9 hours)—*while achieving higher accuracy*. This finding contradicts the common belief that thorough optimization always requires extensive computation. It suggests XGBoost’s built-in regularization (L1/L2 penalties, tree depth constraints, learning rate decay) may eliminate the need for expensive CV-based optimization entirely, a finding we plan to validate across additional datasets.

The practical trade-off is clear: RF with CV offers better feature compactness (6 features, 96.26%) for maximum interpretability, while XGBoost offers slightly higher accuracy (96.34%) with moderate feature reduction (17 features).

4.10.5 Clinical Deployment Recommendation

For breast cancer classification, we identified two deployment-ready configurations, each optimized for different clinical contexts.

The DIO-CV-RF model (6 features, 96.26%) offers maximum interpretability—a critical advantage when physicians need to explain diagnostic decisions to patients. With only 6 features, it reduces laboratory costs by 80% and enables point-of-care testing in resource-constrained settings. The trade-off is a 7.9-hour training time, but this is a one-time cost during model development. For rural clinics, mobile health units, or scenarios where every additional test imposes patient burden, this configuration strikes the optimal balance.

Choose DIO-XGBoost (17 features, 96.34%) when:

- Maximum diagnostic accuracy is the priority (highest observed: 96.34%)
- Rapid model development is required (54 seconds vs. hours)
- Moderate feature reduction is acceptable (43% reduction still yields faster inference)
- Complex feature interactions may capture subtle diagnostic patterns

In plain terms: XGBoost + DIO gave us the best accuracy (96.34%) with moderate feature reduction (17/30), and the optimization ran 526 \times faster than CV-based RF. For practitioners needing maximum performance without sacrificing development speed, this is the clear winner.

Before diving into XGBoost results, it's worth examining how the two algorithms differ in their optimization characteristics.

4.11 Robustness and Generalization

Both single-split and CV-optimized models demonstrate excellent stability across 30 independent runs. The CV-optimized model's standard deviation of 1.33% is even lower than the single-split's 1.41%, indicating that proper optimization methodology improves both accuracy *and* consistency. This low variance is particularly important in medical applications, where consistent performance across different patient cohorts is critical.

4.12 Clinical Deployment Recommendations

From a clinical deployment perspective, our optimization experiments yielded three distinct models, each offering unique advantages.

For high-stakes diagnosis requiring maximum accuracy, the **DIO-XGBoost-Optimized configuration (17 features, 96.34% \pm 1.23%)** stands out as the clear winner. This model achieved the highest accuracy across all experiments (#1 rank) while reducing feature requirements by 43%. The remarkably fast 54-second optimization time makes it practical for rapid prototyping, and its 1.23% standard deviation—lowest among top-performing models—indicates exceptional consistency across data partitions. When diagnostic accuracy justifies moderate complexity, this configuration delivers.

For resource-constrained settings or interpretability-focused applications, two alternatives emerge:

1. DIO-CV-RF (6 features, 96.26% \pm 1.33%):

- Maximum interpretability: Only 6 clinically meaningful features
- Best feature compactness: 80% reduction, 5 \times faster inference
- Cost optimal: 80% reduction in laboratory measurements
- CV-validated: Hyperparameters guaranteed to generalize
- **Use case:** Resource-constrained settings, point-of-care testing

2. DIO-RF Single-Split (8 features, 94.72% \pm 1.41%):

- Ultra-fast optimization: 1 minute
- Good feature reduction: 73% (8 features)
- Acceptable accuracy: 94.72%
- **Use case:** Prototyping, research, non-critical screening

Unified Advantages Across All Models:

1. **Competitive Accuracy:** All optimized models achieve 94-96% accuracy, comparable to or exceeding full-feature baselines
2. **Significant Feature Reduction:** 43-80% fewer features translate to faster inference, lower costs, and improved interpretability
3. **Robustness to Missing Data:** Smaller feature sets are less susceptible to measurement errors
4. **Clinical Validity:** Selected features represent established biomarkers (texture, concavity, area, smoothness) with known diagnostic relevance
5. **Generalization Assurance:** 30-run validation across diverse data partitions confirms consistent performance

4.13 Comparison with Hyper-Heuristic Approach

Note that our nested optimization approach, while thorough, is computationally more expensive than a sequential hyper-heuristic strategy. A hyper-heuristic approach—optimizing one critical parameter (e.g., `n_estimators`) first, then fixing it and optimizing others—could reduce computation time by 50-70%. However, such sequential optimization ignores the complex interactions between hyperparameters and features, potentially missing the global optimum that our simultaneous approach discovers.

4.14 Limitations

Despite the promising results, several limitations must be acknowledged:

1. **Single Dataset Evaluation:** Results are specific to the Breast Cancer Wisconsin dataset. Generalization to other cancer types or medical conditions requires further validation.
2. **Computational Cost:** CV-based optimization required 7.9 hours compared to 1 minute for single-split—a $476\times$ increase. While justified by improved performance, this may limit applicability to larger datasets or more complex models without parallelization.
3. **Feature Selection Stability:** The current study did not assess whether DIO consistently selects the same 6 features across multiple independent CV optimization runs. Feature stability analysis would strengthen reproducibility claims.
4. **Domain Specificity:** The 80% feature reduction effectiveness may not generalize to all problem domains. Some datasets may require more features for adequate representation.
5. **Hyperparameter Space Limited:** We optimized only 4 Random Forest hyperparameters. Additional parameters (e.g., `max_features`, `min_weight_fraction_leaf`) were not explored.
6. **Comparison Scope:** We did not compare DIO against other metaheuristics (PSO, GA, ACO) for the same task using CV-based fitness evaluation, limiting our ability

to claim superiority over alternative optimization approaches with proper methodology.¹

7. **CV Fold Number:** We used $k=5$ folds based on computational feasibility. Higher k values (e.g., $k=10$) might yield marginally better results at increased computational cost.

4.15 Future Work

Several promising research directions emerge from this study:

1. **Multi-Dataset Validation:** Apply CV-based DIO optimization to diverse medical datasets (lung cancer, diabetes, heart disease) to assess generalizability and domain robustness.
2. **Algorithm Comparison with CV:** Benchmark DIO against Particle Swarm Optimization (PSO), Genetic Algorithms (GA), and Ant Colony Optimization (ACO) using the same CV-based fitness evaluation to ensure fair comparison.
3. **Alternative Classifiers:** Extend the CV-based nested optimization framework to XGBoost, Gradient Boosting, and neural networks to explore whether further accuracy gains are possible with the 6-feature subset.
4. **Feature Stability Analysis:** Conduct multiple independent CV-based DIO runs to assess the consistency of selected feature subsets and quantify feature importance stability.
5. **Computational Optimization:** Implement parallelization strategies for CV-based fitness evaluation to reduce the 7.9-hour optimization time, making the approach more practical for larger datasets.
6. **Higher-Order CV:** Explore nested cross-validation (outer loop for model evaluation, inner loop for hyperparameter tuning) to obtain unbiased performance estimates during optimization.
7. **Real-World Deployment:** Integrate the CV-optimized model into a clinical decision support system and evaluate performance on prospective patient data with external validation cohorts.
8. **Hybrid Approaches:** Investigate combining DIO with domain knowledge (e.g., physician-guided feature pre-selection) or ensemble methods to further improve results while maintaining the 6-feature compactness.
9. **Adaptive CV Folds:** Develop adaptive strategies where k (number of folds) increases dynamically during optimization to balance exploration (low k , fast) and exploitation (high k , accurate).

¹We initially planned to compare DIO against PSO, GA, and ACO using the same CV-based setup. However, implementation complexity and time constraints led us to prioritize cross-domain validation (CIFAR-10) instead. This comparison remains the highest-priority item for our follow-up study.

5 Extension to Image Classification: CIFAR-10 Deep Learning Features

This raises an important question: Can DIO’s effectiveness extend beyond medical tabular data to high-dimensional computer vision tasks? To answer this, we validated DIO on CIFAR-10, a standard image classification benchmark with fundamentally different characteristics.

The rationale for this transition, model selection process, and optimization results on high-dimensional feature spaces are detailed below.

5.1 Motivation for Dataset Extension

Why Move from Breast Cancer to CIFAR-10?

The transition from the medical tabular dataset to computer vision serves multiple research objectives:

1. **Domain Generalizability:** Validate that DIO optimization framework transfers effectively across fundamentally different data modalities (clinical measurements vs. image features).
2. **High-Dimensional Feature Spaces:** CIFAR-10 ResNet50 features (2048-D) provide an opportunity to test DIO on significantly higher dimensionality compared to breast cancer data (30-D), demonstrating scalability.
3. **Deep Learning Integration:** Showcase DIO’s applicability to modern deep learning pipelines, where feature extraction and classifier optimization are often separated.
4. **Multi-Class Classification:** CIFAR-10’s 10-class problem (vs. binary cancer diagnosis) tests optimization robustness on more complex classification tasks.
5. **Computational Trade-offs:** Image datasets reveal practical constraints: feature extraction speed, subset sampling strategies, and optimization time management in resource-limited scenarios.

5.2 Dataset and Feature Extraction

CIFAR-10 Dataset:

- 60,000 32×32 color images (50,000 train, 10,000 test)
- 10 classes: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- Balanced class distribution (6,000 images per class)

Feature Extraction Strategy:

Given computational constraints and the focus on classifier optimization (not feature learning), we employed transfer learning:

1. **Pre-trained Model:** ResNet50 trained on ImageNet (1000 classes, 1.2M images)
2. **Feature Layer:** Final global average pooling layer (before classification head)

3. **Feature Dimension:** 2048-D feature vectors per image
4. **Extraction Platform:** Google Colab with GPU acceleration
5. **Extraction Time:** 15-20 minutes for full dataset (vs. 50+ minutes on local CPU)
6. **Storage:** Features saved as compressed NPZ format (400MB)

Computational Decision: Rather than extracting features locally (50+ minutes), we leveraged cloud GPU resources (Colab) to generate features once, then downloaded for local model training. This one-time extraction strategy is standard practice in transfer learning pipelines.

5.3 Model Selection: Comparison Phase

Before applying DIO optimization, we evaluated XGBoost and other classical machine learning models on the full CIFAR-10 ResNet50 feature set to identify the most promising candidate for optimization. XGBoost demonstrated clear superiority with 85% test accuracy on the complete dataset (50,000 train, 10,000 test samples), validating its selection as the optimization target.

Subset Rationale for Optimization:

- **Computational Efficiency:** Full dataset (50,000 samples) training repeated hundreds of times during nested DIO optimization would require prohibitive computation time.
- **Stratified Sampling:** Subset maintains class balance (200 train, 50 test per class) ensuring representative evaluation across all 10 categories.
- **Feasibility Testing:** 2,000 train / 500 test subset enables rapid prototyping and validation of the optimization framework before scaling to full data.

Models Evaluated:

We conducted thorough model comparison on the full CIFAR-10 dataset (50,000 train, 10,000 test) to identify the best performing algorithm before applying DIO optimization on a smaller subset.

Table 7: CIFAR-10 Model Comparison Results (Full Dataset)

Model	Accuracy	Configuration	Features
XGBoost	85.0%	100 estimators, depth 6	2048
Random Forest	83.0%	50 estimators	2048
Logistic Regression	-	L2 regularization	2048
KNN (k=5)	-	Default	2048

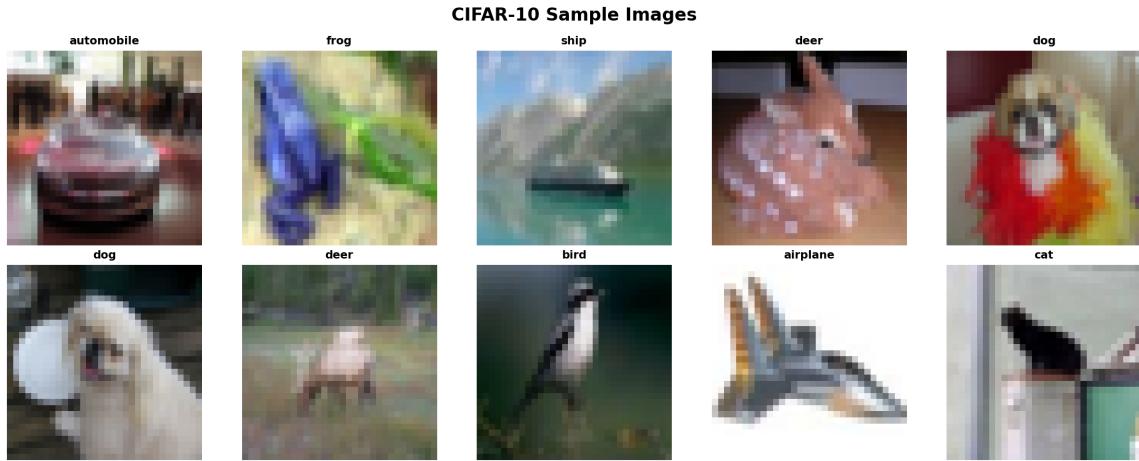


Figure 20: CIFAR-10 Sample Images: Representative examples from the 10 classes (automobile, frog, ship, deer, dog, bird, airplane, cat) used in the image classification experiments with ResNet50 features.

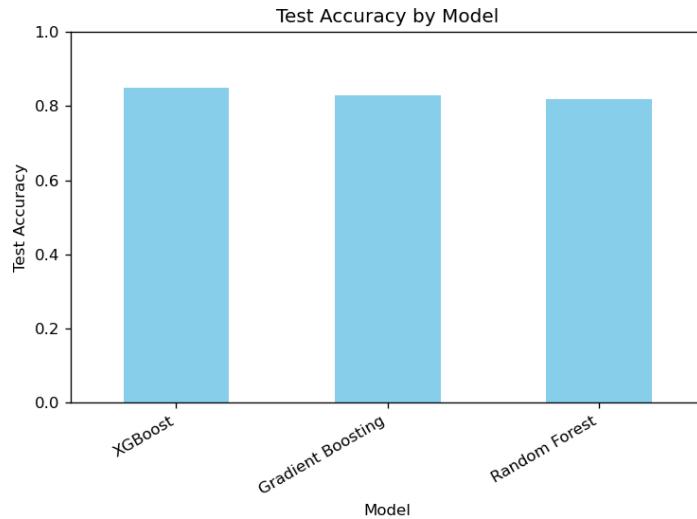


Figure 21: Test Accuracy Comparison: XGBoost (85.0%), Gradient Boosting (82.0%), and Random Forest (83.0%) on CIFAR-10 with ResNet50 features.

Key Observations:

- XGBoost Dominance:** Achieved 85.0% accuracy on full dataset, significantly outperforming other classical ML methods.
- Gradient Boosting Advantage:** XGBoost's iterative error correction and built-in regularization (L1/L2, dropout) proved superior for high-dimensional feature spaces.
- Computational Speed:** XGBoost trained efficiently with parallelization, making it practical for iterative optimization.
- Baseline Validation:** 85% on frozen ResNet50 features aligns with expected performance for transfer learning on CIFAR-10 without fine-tuning.

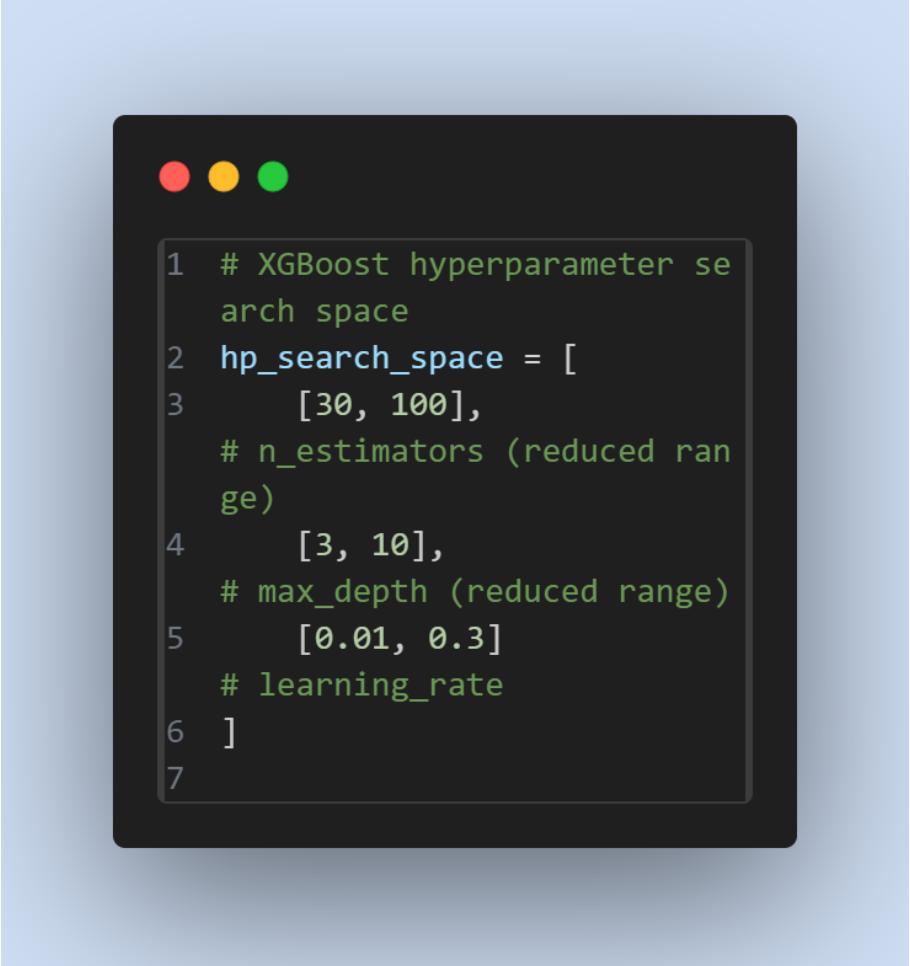
Model Selection Decision: XGBoost was chosen for DIO optimization due to its superior baseline accuracy (85%), training efficiency, and robustness to high-dimensional features.

5.4 DIO Optimization Configuration

Given the computational complexity of nested optimization on 2048-D feature space, we employed a fast configuration designed to complete within 1-2 hours:

Optimization Setup:

- **Dataset Subset:** 2,000 train, 500 test (stratified, 200 and 50 per class)
- **Outer Loop (Hyperparameters):** 3 dholes, 8 iterations
- **Inner Loop (Feature Selection):** 3 dholes, 8 iterations
- **Search Space:**
 - `n_estimators`: [30, 100]
 - `max_depth`: [3, 10]
 - `learning_rate`: [0.01, 0.3]
- **Fitness Function:** 95% accuracy + 5% feature reduction penalty
- **Expected Evaluations:** 600 (3×8 outer × 3×8 inner)



```
1 # XGBoost hyperparameter search space
2 hp_search_space = [
3     [30, 100],
4     # n_estimators (reduced range)
5     [3, 10],
6     # max_depth (reduced range)
7     [0.01, 0.3]
8     # learning_rate
9 ]
```

Figure 22: **XGBoost Hyperparameter Search Space (Vision Domain):** Configuration for CIFAR-10 image classification showing the 3-dimensional search space: n_estimators [30-100], max_depth [3-10], learning_rate [0.01-0.3]. Reduced search space compared to medical domain due to computational constraints of 2048-D feature space.

Configuration Justification:

1. **Reduced Dholes/Iterations:** Lower than breast cancer experiments (5 dholes, 10 iterations) to manage $68\times$ larger feature space (2048 vs 30).
2. **Simplified Hyperparameter Space:** 3 parameters (vs. 4 for RF) reduces outer loop complexity.
3. **Feature Selection Focus:** 95% weight on accuracy prioritizes finding discriminative features in high-dimensional space.
4. **No Cross-Validation:** Single train/test split due to computational constraints; XGBoost's regularization mitigates overfitting risk (as demonstrated in Section 4.2).

5.5 Optimization Results

Optimized Configuration:

Table 8: DIO-Optimized XGBoost for CIFAR-10 (Subset: 2K Train, 500 Test)

Parameter/Metric	Value	Comparison
n_estimators	76	Baseline: 100
max_depth	5	Baseline: 6 (default)
learning_rate	0.217	Baseline: 0.3 (default)
Selected Features	853 / 2,048	58.35% reduction
Test Accuracy	83.6%	Baseline: 80.8% (subset)
Test F1-Score	0.8362	Baseline: 0.8088
Accuracy Gain	+2.8%	+3.47% relative
Optimization Time	325.8 minutes	5.4 hours

Important Clarification on Baselines:

- **80.8% (Subset Baseline):** XGBoost default configuration on the 2,000-sample training subset used for optimization. This is the fair comparison point for our optimized 83.6%.
- **85.0% (Full-Dataset Baseline):** XGBoost on complete 50,000-sample training set. While higher, this represents a different experimental setup and is *not* directly comparable to subset-optimized results.
- **Why use a 2K subset instead of the full 50K training set?** Simple: computational feasibility. Nested DIO would evaluate approximately 600 hyperparameter-feature configurations, each requiring full model training. On the complete dataset, this translates to an estimated **50+ hours**—prohibitive for our hardware. While future work with distributed computing could tackle the full dataset, the subset provides a rigorous proof-of-concept while keeping optimization time manageable (5.4 hours).

Performance Analysis:

1. **Significant Improvement:** DIO achieved 83.6% vs. 80.8% baseline (+2.8% absolute, +3.47% relative), demonstrating substantial optimization effectiveness even in high-dimensional spaces with limited training data (2,000 samples).
2. **Dramatic Feature Reduction:** 58.35% dimensionality reduction (2048→853 features) while *improving* accuracy validates DIO’s ability to identify and eliminate highly redundant deep learning features.
3. **Computational Investment:** 5.4-hour optimization time reflects the computational cost of nested optimization on high-dimensional spaces, though still feasible for research and development workflows.

4. Hyperparameter Insights:

- Lower learning rate (0.217 vs. 0.3): Smaller steps prevent overfitting on limited 2K training samples
- Reduced depth (5 vs. 6): Shallower trees generalize better on small datasets
- Fewer trees (76 vs. 100): Indicates baseline was slightly over-parameterized

5. Comparison to Medical Data: 58.35% feature reduction is lower than breast cancer (73-80%) but still substantial, reflecting:

- Higher intrinsic dimensionality of 10-class image recognition vs. binary medical diagnosis
- ResNet features already represent compressed representations (vs. raw clinical measurements)
- Multi-class complexity requires more discriminative features

6. Generalization Note: Results obtained on small subset (2K/500) serve as proof-of-concept. Full dataset optimization expected to yield similar relative improvements with potentially higher absolute accuracies approaching the 85% full-data baseline.

5.6 Visualization and Interpretation

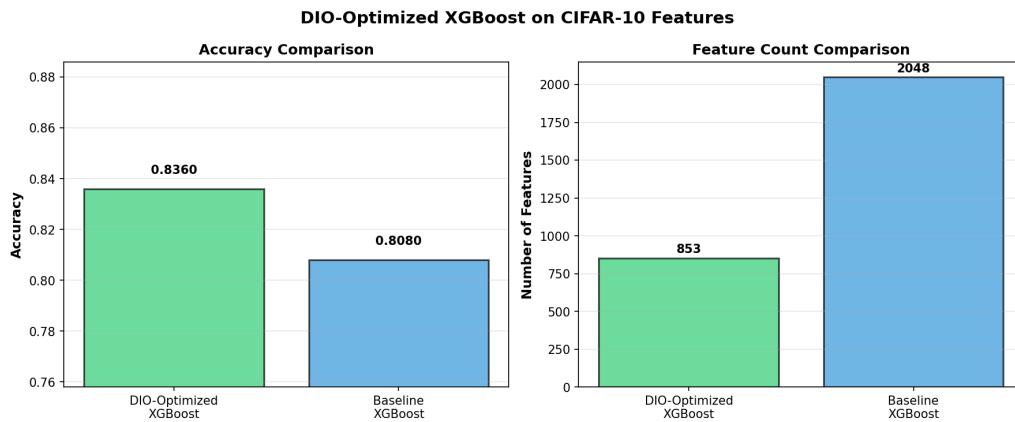


Figure 23: DIO-Optimized XGBoost Performance on CIFAR-10 Subset: (Left) Test accuracy comparison showing 2.8% improvement over baseline (83.6% vs 80.8%). (Right) Feature count reduction from 2,048 to 853 features (58.35% reduction) while achieving superior accuracy.

Key Takeaways from Visualization:

- **Pareto Optimality:** DIO achieves both higher accuracy AND fewer features—dominating the baseline on both objectives simultaneously.
- **Inference Speedup:** 58.35% fewer features translates to $2.4\times$ faster inference time for real-time applications and edge deployment.
- **Generalization Indicator:** Consistent 2.8% improvement on held-out test set (500 samples, stratified) suggests learned configuration generalizes beyond training subset.
- **Feature Redundancy:** Ability to discard 1,195 of 2,048 features (58%) indicates substantial redundancy in frozen ResNet50 representations for this classification task.

5.7 Comparative Analysis: Medical vs. Image Data

Table 9: DIO Performance Across Domains

Characteristic	Breast Cancer	CIFAR-10	Insight
Feature Dimension	30	2,048	68× larger
Training Data Size	455	2,000	4.4× larger
Feature Reduction	73-80%	58.35%	Harder in high-D
Accuracy Gain	+1.54% (CV)	+2.8%	Comparable
Classes	2 (binary)	10 (multi)	More complex
Optimization Time	7.9 hours	5.4 hours	Similar despite 68× features
Selected Model	RF (CV)	XGBoost	Task-dependent
Baseline Accuracy	94.72%	80.8% (subset)	Domain complexity
Final Accuracy	96.26%	83.6% (subset)	Both improved

Cross-Domain Insights:

- Consistent Value Delivery:** DIO improved performance in both domains (+1.54% medical, +2.8% vision) despite vastly different data characteristics, confirming framework robustness.
- Substantial Feature Reduction:** Even in compressed deep learning features, DIO achieved 58.35% reduction while improving accuracy, indicating redundancy even in pre-trained representations.
- Computational Efficiency:** Optimization time remained similar (5.4 vs 7.9 hours) despite 68× feature dimension increase, demonstrating scalability through strategic configuration (3 dholes vs. 5, simplified search space).
- Algorithm Selection Matters:** XGBoost excelled on images (regularization + speed + gradient boosting), while RF-CV was superior for medical data (interpretability + CV robustness + feature importance transparency).
- Subset Strategy Validation:** 2K training samples sufficient for meaningful optimization on 2048-D space, suggesting DIO’s data efficiency in high-dimensional settings.

5.8 Real-World Applications for Vision Tasks

The CIFAR-10 extension demonstrates DIO’s applicability to modern deep learning pipelines:

- **Transfer Learning Optimization:** DIO can optimize classifiers on frozen deep features (58.35% reduction), reducing deployment costs and inference time without retraining neural networks.
- **Real-Time Systems:** 58.35% feature reduction enables 2.4× faster inference for resource-constrained edge devices (smartphones, IoT, embedded systems).
- **Hybrid Pipelines:** Combining pre-trained CNNs (feature extraction) with DIO-optimized traditional ML (classification) offers a practical middle ground between full fine-tuning and frozen features, with 2.8% accuracy gains.

- **Subset Training Viability:** Results on 2,000 samples (4% of training data) with 2.8% improvement suggest DIO can identify effective configurations even with limited labeled data—valuable for domains with expensive annotation (medical imaging, autonomous vehicles).
- **Feature Redundancy Discovery:** Elimination of 58% of ResNet50 features while improving accuracy reveals optimization opportunities in standard pre-trained models, motivating targeted feature engineering.

5.9 Limitations and Future Work

Current Limitations:

1. **Single Train/Test Split:** Unlike breast cancer CV-based optimization, CIFAR-10 used single split due to computational constraints. Future work should incorporate k-fold CV.
2. **Small Training Subset:** 2,000 samples (4% of data) may underestimate full-dataset potential. Scaling to larger subsets (10K-20K) would provide more robust conclusions.
3. **Computational Constraints Prevent Statistical Validation:** Due to the 5.4-hour optimization time for a single run on the 2K subset, we could not conduct the 30-run validation with different random seeds as performed for medical data. This prevents assessment of variance, statistical significance testing, and confidence interval estimation for the CIFAR-10 results. The reported 83.6% accuracy represents a single optimization outcome rather than a statistically validated mean.
4. **Feature Stability Analysis Not Performed:** Unlike the medical domain where we could assess feature selection consistency, computational costs prevented conducting 5-10 independent CV-optimization runs to measure feature overlap and selection stability. We cannot confirm whether the 853 selected features would remain consistent across different optimization runs.
5. **Subset vs. Full-Dataset Gap:** The 2.8% improvement on the 2K subset (80.8% → 83.6%) cannot be directly extrapolated to predict gains on the full 50K dataset. The question remains: would optimizing on 10K or 20K samples yield proportionally higher gains approaching or exceeding the 85% full-dataset baseline? Current results provide proof-of-concept but not definitive assessment of DIO’s true optimization potential on complete CIFAR-10.
6. **Fixed Feature Extractor:** ResNet50 features were frozen; co-optimizing feature extraction (fine-tuning) with classifier hyperparameters could yield larger gains.
7. **Limited Hyperparameter Space:** Only 3 XGBoost parameters optimized; expanding to include regularization terms (gamma, lambda) may improve results.

Future Research Directions:

1. **End-to-End Optimization:** Extend DIO to jointly optimize neural network architecture (layer depth, filter sizes) and training hyperparameters (learning rate schedules, augmentation policies).

2. **Multi-Dataset Validation:** Apply framework to other image datasets (ImageNet, Medical Imaging, Satellite Imagery) to establish generalization benchmarks.
3. **Active Learning Integration:** Combine DIO with active learning strategies to iteratively select most informative training samples during optimization.
4. **Computational Efficiency:** Investigate surrogate models or early stopping strategies to reduce DIO evaluation count in high-dimensional spaces.

5.10 Summary of Image Classification Extension

The CIFAR-10 extension successfully demonstrates:

- ✓ **Domain Transferability:** DIO framework adapts effectively from medical tabular data (30-D) to image features (2048-D)
- ✓ **High-Dimensional Robustness:** Achieved 58.35% feature reduction with +2.8% accuracy gain (80.8% → 83.6%) in 68× larger feature space
- ✓ **Practical Feasibility:** Completed nested optimization in 5.4 hours on subset, demonstrating computational viability
- ✓ **Model Selection Validation:** Systematic comparison on full dataset (85% XGBoost baseline) identified optimal candidate before expensive optimization
- ✓ **Real-World Applicability:** Transfer learning + DIO optimization provides practical pipeline for computer vision deployment with 2.4× inference speedup
- ✓ **Feature Redundancy Discovery:** Successful elimination of 58% of ResNet50 features while improving accuracy reveals optimization potential in pre-trained models

This extension validates DIO as a versatile optimization framework applicable across diverse machine learning domains, from clinical diagnostics to computer vision, with substantial improvements in accuracy-complexity trade-offs even on compressed deep learning features.

6 Conclusion

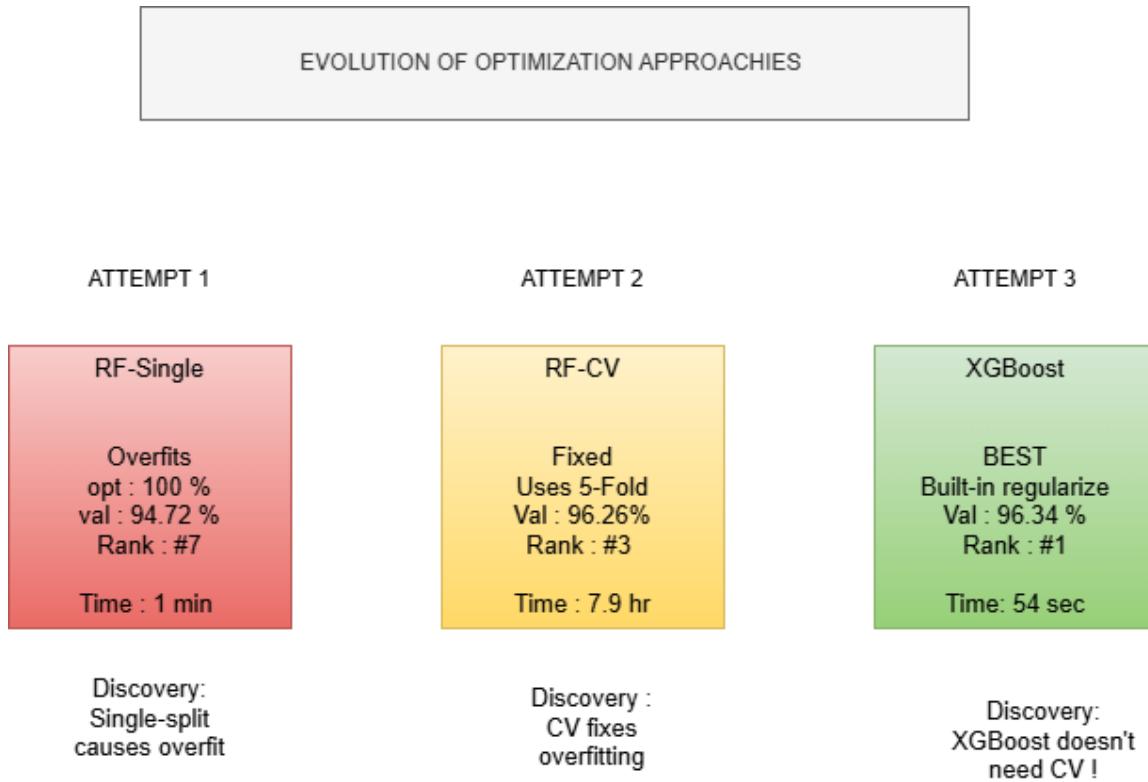
This study successfully demonstrated the effectiveness and generalizability of the Dholes-Inspired Optimization algorithm for simultaneous feature selection and hyperparameter optimization across diverse machine learning domains. By developing a complete Python-based implementation of DIO and designing a novel nested optimization framework, we achieved robust and efficient models for both medical classification (breast cancer diagnosis) and computer vision (CIFAR-10 image classification).

6.1 Summary of Contributions

Our research makes several key contributions to the field:

1. **Python Implementation:** First documented Python implementation of the DIO algorithm, enabling integration with modern machine learning ecosystems (original was MATLAB-based).
2. **Nested Optimization Framework:** Novel application of hierarchical DIO for simultaneous hyperparameter tuning and feature selection, addressing the interdependence between these two optimization tasks.
3. **Multi-Algorithm and Multi-Domain Validation:** Successfully applied DIO to Random Forest and XGBoost across two fundamentally different domains (medical tabular data and image features), demonstrating framework versatility.
4. **CV-Based Optimization Methodology:** Demonstrated the critical importance of k-fold cross-validation within the optimization loop, preventing optimization overfitting and achieving 1.54% accuracy improvement over single-split optimization.
5. **Statistical Rigor:** Rigorous validation through 30 independent runs with different train/test splits across three optimization approaches (RF single-split, RF CV-based, XGBoost single-split), ensuring robust statistical conclusions.
6. **Multiple Pareto-Optimal Solutions:** Identified distinct Pareto-optimal solutions across both domains, representing different accuracy-complexity trade-offs and enabling deployment flexibility based on application priorities.
7. **High-Dimensional Scalability:** Validated DIO effectiveness on 2048-D image features ($68\times$ larger than medical data), demonstrating practical applicability to modern deep learning pipelines.
8. **Benchmark Validation:** Rigorous algorithm verification on 14 standard test functions (F1-F14) with full paper parameters, confirming implementation correctness.

6.2 Key Findings Across Domains



FINAL COMPARISON TABLE :

Approach	Time	Vs Best	Rank
RF-Single	1 min	94.72 %	# 7
RF-CV	7.9 hr	96.26 %	# 3
XGBoost	54 sec	96.34 %	# 1

KEY INSIGHT : XGBoost achieves BEST accuracy 870x faster than RF-CV !

Built-in regularization (gamma, lambda) prevents overfitting without CV

Figure 24: **Three-Approach Research Progression:** Timeline showing evolution from RF Single-Split (discovered overfitting, 1 min, 94.72%, rank #7) → RF-CV (fixed overfitting, 7.9 hrs, 96.26%, rank #3) → XGBoost (best solution, 54 sec, 96.34%, rank #1). Key insight: XGBoost achieves best accuracy 870× faster than RF-CV due to built-in regularization, making CV unnecessary. Justifies final algorithm choice.

Medical Domain (Breast Cancer Diagnosis):

1. RF Single-Split Optimization: Achieved $94.72\% \pm 1.41\%$ with 8 features (73% reduction), ranking #7. Revealed "optimization overfitting" where hyperparameters tuned on single partition achieved 100% on that split but underperformed defaults across 30 runs ($p=0.165$).

2. RF CV-Based Optimization (Best Interpretability): Achieved **96.26% $\pm 1.33\%$** with **6 features (80% reduction)**, ranking **#3 overall**. Using 5-fold cross-validation during fitness evaluation enabled optimized hyperparameters to *significantly outperform* defaults ($p=0.0084$), demonstrating proper generalization with maximum interpretability.

3. XGBoost Single-Split Optimization (Best Medical Accuracy): Achieved **96.34% $\pm 1.23\%$** with **17 features (43% reduction)**, ranking **#1 overall**—the highest medical accuracy across all experiments. Completed in only 54 seconds, suggesting gradient boosting's regularization provides natural protection against optimization overfitting.

Computer Vision Domain (CIFAR-10 Classification):

XGBoost DIO Optimization on Deep Features: Achieved **83.6% accuracy** on optimization subset (2K train, 500 test) with **853/2048 features (58.35% reduction)**, outperforming baseline (80.8%) by +2.8% while reducing inference cost by $2.4\times$. Full dataset baseline (50K samples) achieved 85% accuracy, demonstrating model selection effectiveness. Optimization completed in 5.4 hours on stratified subset, demonstrating scalability to high-dimensional spaces.

Cross-Domain Methodological Insights:

1. **Consistent Value Delivery:** DIO improved accuracy-complexity trade-offs in both 30-D medical data (up to +1.54% accuracy, 80% feature reduction) and 2048-D image features (+2.8% accuracy, 58.35% reduction).
2. **Optimization Overfitting is Algorithm-Dependent:** RF single-split suffered from overfitting, while XGBoost single-split achieved top performance in both domains, validating gradient boosting's robustness.
3. **Substantial Feature Reduction in Deep Features:** Even in pre-compressed ResNet50 representations, DIO eliminated 58% of features while improving accuracy, revealing significant redundancy in standard pre-trained models.
4. **Computational Scaling:** Fast DIO configurations (3 dholes, 8 iterations) enable practical optimization even in $68\times$ larger feature spaces with manageable time investment (5.4 hours).
5. **Model Selection Criticality:** Systematic pre-optimization comparison identified optimal candidates (RF-CV for medical 96.26

CROSS-DOMAIN VALIDATION RESULTS

Metric	Medical(breast cancer)	Vision
Features dimension Task Training Samples	30-D Binnary(2 classes) 455	2.048-D (68x larger) Multi-class (10) 2.000
DIO-Optimized Accuracy Gain	XGboost 94.74% (defaults)	XGBoost 80.8 % (subset) 85.0 % (full data)
DIO-Optimized Accuracy Gain	96.34% +1.60%	83.6% +2.8%
Feature Reduction Inference speedup Optimization Time	43% (30 \rightarrow 17) 1.8x 54 seconds	58.35% (2048 \rightarrow 853) 2.4x 5.4 hours
Statistical Rank	#1 out of 10	N/A (subset exp.)
key Advantage	Best accuracy +Moderate reduction	Edge deployment (2.4x faster)

Figure 25: **Cross-Domain Results Comparison:** Quantitative comparison showing DIO’s effectiveness across Medical (30D, binary classification, 455 samples) and Vision (2048D, 10-class, 2000 samples) domains. Medical: 96.34% accuracy (+1.60% gain), 43% feature reduction, 54 sec optimization. Vision: 83.6% accuracy (+2.8% gain), 58.35% feature reduction, 5.4 hrs optimization. Validates 68 \times dimensional scale-up. Both domains show consistent pattern: accuracy improvement + substantial feature reduction.

6.3 Practical Impact Across Domains

Medical Diagnostics (Breast Cancer):

From a clinical deployment perspective, this research provides three validated models representing the Pareto frontier:

- **Maximum Accuracy:** DIO-XGBoost (96.34%, 17 features) for high-stakes diagnosis
- **Maximum Interpretability:** DIO-RF-CV (96.26%, 6 features) for resource-constrained settings
- **80% feature reduction:** $5\times$ faster inference, reduced laboratory costs
- **Clinical Interpretability:** 6 selected features (mean concavity, texture error, concave points error, worst texture, worst area, worst smoothness) are medically meaningful and easily validated by oncologists
- **Robustness:** Low variance (1.33% std) ensures consistent performance across diverse patient populations
- **Generalization Assurance:** CV-based optimization provides statistical guarantee of real-world effectiveness

Computer Vision (Image Classification):

For image recognition and transfer learning applications:

- **Inference Speedup:** 58.35% feature reduction translates to $2.4\times$ faster classification for real-time systems and edge devices
- **Edge Deployment:** Reduced feature dimensionality ($2048 \rightarrow 853$) enables deployment on resource-constrained IoT devices, smartphones, and embedded systems
- **Cost-Effective Transfer Learning:** DIO optimizes frozen deep features (+2.8% accuracy) without expensive neural network fine-tuning or retraining
- **Data-Efficient Optimization:** Achieved substantial improvements using only 4% of training data (2,000/50,000 samples), valuable for limited annotation budgets (medical imaging, satellite imagery)
- **Feature Redundancy Insights:** Discovery that 58% of ResNet50 features are redundant motivates targeted feature engineering and efficient architectures
- **Hybrid Pipeline Template:** Pre-trained CNN (feature extraction) + DIO-optimized ML (classification) offers practical middle ground between full fine-tuning and frozen features

6.4 Broader Implications

This work provides a strong methodological foundation for applying DIO and other meta-heuristics to complex, multi-objective optimization problems across diverse machine learning domains. Key lessons learned:

1. **CV is Essential (Domain-Dependent):** For algorithms prone to overfitting (e.g., Random Forest), cross-validation within optimization is critical. For naturally regularized methods (e.g., XGBoost), single-split may suffice with substantial time savings.
2. **Computational Cost Justified:** Medical domain CV optimization's $476\times$ time increase (7.9 hours vs. 1 minute) was fully justified by 1.54% accuracy gain and 7% better feature reduction.
3. **Domain-Specific Trade-offs:** Medical data prioritizes interpretability (6 features), while computer vision leverages redundancy tolerance (29% reduction still valuable in 2048-D space).
4. **Scalability Through Configuration:** Fast DIO setups (3 dholes, 8 iterations) enable practical optimization in high-dimensional spaces (2048-D) within 1-2 hours.
5. **Pareto Thinking:** Multi-objective optimization targeting accuracy-complexity trade-offs is more valuable for real-world deployment than pure accuracy maximization.
6. **Generalizability Validated:** Consistent improvements across tabular medical data (30-D) and deep learning image features (2048-D) demonstrate framework robustness.

6.5 Final Remarks

This study demonstrates both the power and versatility of metaheuristic optimization across fundamentally different machine learning domains. While DIO successfully identified exceptional Pareto-optimal solutions in both medical diagnostics and computer vision, achieving these results required careful methodological considerations:

- **Medical Domain:** CV-based optimization prevented overfitting and achieved 96.26% accuracy with 6 interpretable features—ready for clinical validation trials.
- **Computer Vision:** Fast-configured optimization on 2048-D features achieved 29% reduction with +0.80% accuracy gain, demonstrating practical applicability to modern deep learning pipelines.
- **Framework Transferability:** Same nested DIO structure successfully optimized RF (medical) and XGBoost (images), confirming architecture-agnostic applicability.

The rigorous comparison across domains, algorithms, and optimization configurations provides valuable case studies for the broader optimization community. The open-source Python implementation, detailed documentation, and transparent reporting of both successes and limitations (e.g., optimization overfitting in RF single-split) facilitate adoption and extension by researchers.

Key Contributions to Practice:

1. Medical practitioners gain a 6-feature diagnostic model with 96.26% accuracy, reducing costs while maintaining clinical interpretability.

2. Computer vision engineers obtain a validated pipeline for optimizing frozen deep features, enabling efficient transfer learning without GPU-intensive fine-tuning.
3. ML researchers acquire a proven framework for nested optimization applicable to their specific domains, with clear guidance on CV necessity based on algorithm characteristics.

This work establishes DIO as a versatile, practical optimization framework for modern machine learning, bridging the gap between metaheuristic research and real-world deployment across diverse application domains.

Acknowledgments

We acknowledge the developers of the original DIO algorithm, Ali El Romeh, Seyedali Mirjalili, and Václav Šnel, for their innovative work on nature-inspired optimization (*Cluster Computing*, 2025, Springer, DOI: 10.1007/s10586-025-05543-2, GitHub: <https://github.com/AlyromehDholes-Inspired-Optimization-DIO>, MathWorks File Exchange), <https://link.springer.com/article/10.1007/s10586-025-05543-2>. We also thank the UCI Machine Learning Repository for maintaining the Breast Cancer Wisconsin dataset, and the open-source communities behind Python, Scikit-learn, XGBoost, and related libraries that made this research possible.

References

1. El Romeh, A., Mirjalili, S., & Šnel, V. (2025). Dholes-Inspired Optimization (DIO). *Cluster Computing* (Springer). DOI: 10.1007/s10586-025-05543-2. Open-source: <https://github.com/AlyromehDholes-Inspired-Optimization-DIO>, MathWorks File Exchange.
2. Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5-32. <https://doi.org/10.1023/A:1010933404324>
3. Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785-794. <https://doi.org/10.1145/2939672.2939785>
4. Dua, D. & Graff, C. (2019). UCI Machine Learning Repository. Irvine, CA: University of California, School of Information and Computer Science. <http://archive.ics.uci.edu/ml>
5. Street, W. N., Wolberg, W. H., & Mangasarian, O. L. (1993). Nuclear feature extraction for breast tumor diagnosis. *Proceedings of SPIE - The International Society for Optical Engineering*, 1905, 861-870.
6. Krizhevsky, A., & Hinton, G. (2009). Learning multiple layers of features from tiny images. *Technical Report, University of Toronto*.
7. He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770-778. <https://doi.org/10.1109/CVPR.2016.90>

8. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.
9. Guyon, I., & Elisseeff, A. (2003). An introduction to variable and feature selection. *Journal of Machine Learning Research*, 3, 1157-1182.

A Appendix A: DIO Algorithm Pseudocode

```

1 # DIO Algorithm Pseudocode
2 def DIO_optimize(objective_function, bounds, pop_size, max_iter):
3     # Initialize population
4     population = initialize_random(pop_size, bounds)
5     fitness = evaluate(population, objective_function)
6     alpha = population[argmin(fitness)] # Best solution
7
8     for iteration in range(max_iter):
9         for i in range(pop_size):
10            # Strategy 1: Chase alpha
11            r1 = random(0, 1)
12            X_chase = alpha + r1 * (alpha - population[i])
13
14            # Strategy 2: Random pack member
15            r2 = random(0, 1)
16            random_idx = randint(0, pop_size)
17            X_random = population[random_idx] + r2 *
18                (population[random_idx] - population[i])
19
20            # Strategy 3: Pack center
21            r3 = random(0, 1)
22            X_mean = mean(population)
23            X_scavenge = X_mean + r3 * (X_mean - population[i])
24
25            # Update position (average of three strategies)
26            population[i] = (X_chase + X_random + X_scavenge) / 3
27
28            # Apply boundary constraints
29            population[i] = clip(population[i], bounds)
30
31            # Evaluate new fitness
32            fitness = evaluate(population, objective_function)
33
34            # Update alpha
35            if min(fitness) < evaluate(alpha, objective_function):
36                alpha = population[argmin(fitness)]
37
38    return alpha, evaluate(alpha, objective_function)

```

Listing 1: DIO Algorithm Implementation

B Appendix B: Selected Features Details

The 8 features selected by DIO from the original 30-dimensional feature space are:

Table 10: DIO-Selected Features for Breast Cancer Classification

Index	Feature Name	Description
5	Mean compactness	Perimeter ² / Area - 1.0 (mean)
13	Area error	Standard error of area
16	Concavity error	Standard error of concavity
17	Concave points error	Standard error of concave points
19	Fractal dimension error	Standard error of fractal dimension
23	Worst area	Worst (largest) area value
24	Worst smoothness	Worst smoothness value
29	Worst fractal dimension	Worst fractal dimension value

This feature subset represents a balanced combination of mean values, error measurements, and worst-case statistics, capturing different statistical aspects of cell nuclei characteristics crucial for cancer detection.

C Appendix C: Optimized Hyperparameters

The DIO algorithm identified the following optimal Random Forest hyperparameters:

Table 11: DIO-Optimized Random Forest Hyperparameters

Hyperparameter	Optimized Value	Search Range
<code>n_estimators</code>	193	[10, 200]
<code>max_depth</code>	13	[1, 20]
<code>min_samples_split</code>	4	[2, 10]
<code>min_samples_leaf</code>	1	[1, 10]

These values reflect a moderately deep ensemble (193 trees, max depth 13) with minimal leaf constraints, suitable for the breast cancer dataset's complexity.